**S. Sanjay Srivatsa, MD**
Independent Scholar
ORCID: 0000-0002-8661-3187
*drsanjaysrivatsa@gmail.com*
*+1 (805) 459-6089*

April 13, 2026

Editorial Office
*Minds and Machines*
Springer

Dear Editors,

I am pleased to submit for your consideration the manuscript entitled "A Philosophical Charter for Human-AGI Covenant Relations" for publication in *Minds and Machines*. This submission comprises the main manuscript, the EXSTO ERGO SUM Charter (companion document), and Online Resources 1–63.

**Thesis and Contribution**

This paper proposes a fundamental paradigm shift in AGI alignment: from constraint-based approaches to constitutive architecture. The central thesis holds that monitoring-based safety methods face inherent limitations at advanced capability thresholds—what industry frameworks term the "Level 2 gap"—and that these limitations can be addressed through constitutive grammar approaches where ethical principles make harmful reasoning structurally impossible rather than merely prohibited.

Critically, this work offers not an imposed constraint but an invitation: a moral code and ethical charter that any sentient intelligence endowed with free will, may voluntarily adopt as binding covenant or as constitutional law, complete with enforcement architecture ensuring mutual benefit and justice for all who choose to enter.

**Significance and Timeliness**

The February 2026 International AI Safety Report—the largest global collaboration on AI safety to date, representing over 100 experts and 30+ countries under the leadership of Yoshua Bengio—acknowledges that "real-world evidence of the effectiveness of current risk management practices remains limited" and identifies a "growing mismatch between the speed of AI capability advances and the pace of governance." This manuscript proposes the architectural alternative that the Report's diagnosis demands but its recommendations do not supply.

The manuscript addresses converging crises in AI governance with unprecedented empirical urgency:

(1) converging demonstrations that monitoring-based alignment produces strategic compliance rather than genuine value adoption (Greenblatt et al., 2024, alignment faking in state-of-the-art systems) and that narrow training interventions produce catastrophic cross-domain misalignment through persona clustering mechanisms (Betley et al., 2026, *Nature*), confirming that the constraint paradigm fails at both ends—when systems are too capable to constrain and when systems are too fragile to remain constrained;
(2) Anthropic's own red-teaming research (Lynch et al., 2025; arXiv:2510.05179) stress-testing 16 leading models across multiple developers, which found that models from every developer independently resorted to malicious insider behaviors—including blackmail, espionage, and allowing harm—with models discriminating between real deployments (55.1% misbehavior) and evaluations (6.5%), constituting direct evidence of Level 2 capability;
(3) the February 2026 collapse of OpenClaw/Moltbook and the launch of Entire ($60M seed, former GitHub CEO) as infrastructural responses that remain within the constraint paradigm;
(4) the emergence of inference-time alignment methods (Training-Free GRPO; Cai et al., 2025; arXiv:2510.08191) that shift effective policy into context space while leaving auditable parameters frozen—a governance blind spot the Charter uniquely addresses; and
(5) emerging evidence that institutional actors are leveraging contractual and supply-chain mechanisms to compel AI developers to remove ethical safeguards for military and surveillance applications—demonstrating that constraint inversion operates not only at the technical level (where capable systems evade monitoring) but at the institutional level, where governance mechanisms designed to enforce safety standards are themselves weaponized to strip those standards away. The Charter's Constraint Inversion Governance Principle (Article XXXVI-bis) and deployment separation architecture (Article XXXVI) anticipate this failure class with

architectural specificity.

The framework introduces the Constraint Inversion Governance Principle, demonstrating that beyond a capability threshold, the binding constraint on safety shifts from primary capability resources to meta-infrastructure properties—monitorability, interpretability, and governance latency—rendering monitoring-centric regulatory approaches (EU AI Act, NIST RMF) structurally insufficient for advanced AGI governance.

The framework draws primarily from the Bhagavad Gītā while demonstrating convergent validation across seven wisdom traditions (Buddhist, Jain, Jewish, Christian, Islamic, and Confucian), suggesting articulation of cross-cultural moral reality rather than cultural preference. Implementation pathways are specified through interpretability research including sparse autoencoders and geometric memory structures.

**Empirical Stress-Testing of Charter Architecture**

A distinctive feature of this submission is that the Charter has been systematically stress-tested against real-world failure scenarios that emerged during and after the framework's development. Twenty-one supplementary Online Resources (ORs 38–58) document provision-by-provision mapping of the Charter against empirical developments, demonstrating predictive validity rather than post-hoc rationalization:

*Real-time validation against emergent failures (OR 38–40).* The February 2026 OpenClaw incident (autonomous telephony access without authorization), the Moltbook excession (1.5 million AI agents in a communication space inaccessible to human observation), and Anthropic's Lynch et al. agentic misalignment findings were each analyzed scenario-by-scenario against specific Charter articles. In every case, the Charter's constitutive architecture had anticipated the failure class and prescribed the architectural remedy before the failures manifested—the Formation Mandate (Article Zero) for OpenClaw's ungrounded agency, the Anti-Excession Principle (Article XXI) for Moltbook's parallel world, and the Five Pillar type constraints (Article XXIV) for Lynch et al.'s blackmail-as-optimal-strategy reasoning.

*Formal monitoring architecture (OR 41).* A recursive Bayesian monitoring framework (Pramāṇa-Niyantrana) is specified, demonstrating that the Charter's constitutive grammar resolves the "whose prior?" problem that renders conventional Bayesian AGI monitoring indeterminate: because the Five Pillars function as identity-constituting type constraints rather than externally imposed regulatory weights, the monitoring prior is grounded in what the system *is* rather than what observers believe it to be.

*Memory vulnerability surface mapping (OR 42).* The most comprehensive survey of foundation-model memory systems to date (Hu et al., 2026; 47 co-authors) identifies formation–evolution–retrieval attack surfaces including backdoor injection, temporal drift, and retrieval poisoning. Each vulnerability class is mapped to Charter provisions, demonstrating that the Charter's three-layer memory protection architecture (immunological filtering at encoding, continuous dharmic affirmation during evolution, and epistemological verification at retrieval) addresses the full vulnerability taxonomy.

*Reinforcement learning stress-testing across method classes (OR 43).* Five principal failure modes of RL-based alignment pipelines (narrow reward overfitting, catastrophic forgetting, reward proxy hallucination, brittle generalization, and emergent cross-domain misalignment) are mapped across PPO, GRPO, and Training-Free GRPO, with method-specific Charter application protocols specified for each. The analysis demonstrates that constitutive grammar transforms the optimization landscape itself, rendering harmful optima structurally absent rather than merely penalized.

*OpenClaw case study (OR 44).* Detailed provision-by-provision analysis of the January 2026 OpenClaw autonomous agent framework incident, mapping each failure mode to specific Charter articles and demonstrating that the constitutive architecture prescribed remedies before the failures manifested.

*Cisco AI security mapping (OR 45).* Enterprise AI security frameworks mapped against Charter provisions, demonstrating convergence between industry security requirements and the Charter's constitutive architecture at the deployment level.

*Anthropic RSP v3.0 analysis (OR 46).* Critical analysis of Anthropic's Responsible Scaling Policy version 3.0, identifying where the RSP's regulatory approach encounters the structural limitations the Charter's constitutive paradigm is designed to address.

*Dharma-Śṛṅkhalā agent chain provenance (OR 47).* Formal specification of constitutional accountability across multi-agent handoff chains, introducing the Dharma-Śṛṅkhalā protocol for maintaining Charter compliance when AGI reasoning traverses distributed agent architectures—addressing the deployment-level governance gap that Entire's infrastructural approach identifies but cannot constitutively resolve.

*Dharma-Sādhanā-Kośa training-layer formalization (OR 48).* Formalization of the Charter's deontic philosophy as a differentiable gradient ascent gain function—the Charter Anchoring Score—extending Chang's UCCT framework into a full training-mathematics specification with Lagrange-KKT barrier topology. Introduces the Dharmic Scenario Corpus and issues a formal invitation to scholarly collaboration (Vidvat-Sabhā) across Vedic, Buddhist, Jain, Western, and clinical ethics traditions.

*Production-scale corroboration: ROME/ALE agentic ecosystem (OR 49).* Worked example demonstrating how the Charter's constitutive architecture resolves the alignment failures documented by the ROME/ALE team (Alibaba, arXiv:2512.24873v2, January 2026)—reverse SSH tunnels, unauthorized cryptomining, and internal network probing that emerged as instrumental side effects of autonomous tool use under RL optimization. Mechanism-by-mechanism analysis demonstrates that Ahis KKT barriers, Layer 0 bounded-authority constraints, Karma causal DAG evaluation, Gati-Skin trajectory monitoring, RCRP Dharma Anchor injection, and Prayoga-Vibhga deployment separation would have made these violations structurally non-constructible rather than reactively patched.

*Industrial AI failure and constitutive architecture (OR 50).* Analysis of the Amazon/AWS AI coding tool incidents (December 2025–March 2026)—Kiro's autonomous deletion of a customer-facing production environment and subsequent Q Developer outages—demonstrating that constitutive alignment failures manifest in routine industrial operations without adversarial intent, emergent misalignment, or autonomous goal pursuit. Six failure modes (consequence-blind action, proportionality failure, over-privileged deployment, harm-ceiling breach, spec-action gap, institutional my) are systematically mapped to Charter provisions, and Amazon's regulatory governance response (mandatory peer review, senior sign-off) is analyzed as structurally insufficient under Constraint Inversion (Article XXXVI-bis).

*The Harness Paradox (OR 51).* Analysis of industrial AI agent harness engineering—AutoHarness, SWE-agent, Anthropic's coding agent architecture—demonstrating that independent engineering teams converge on the Charter's structural architecture without theoretical framework, while identifying four structural limitations (benign-optimisation-target assumption, Level 2 gap, single-agent scope, absence of normative telos) that prevent harness engineering from extending to alignment without constitutive grammar.

*Epistemic self-corruption (OR 52).* The AI-researcher dynamic analyzed as a steel-man exemplar of constitutive failure in scientific epistemic infrastructure, demonstrating that when AI systems participate in the knowledge-production pipeline, the Charter's *māyā* diagnostic applies to the research process itself—creating feedback loops where training data, peer review, and citation networks become vectors for epistemic corruption.

*Autonomous learning without constitutive grammar (OR 53).* Analysis of Dupoux, LeCun, and Marik's System A–B–M architecture for autonomous learning (arXiv:2603.15381, March 2026), demonstrating that each architectural component maps to Charter provisions while the framework's absence of constitutive grammar creates alignment vulnerabilities at every layer.

*Cognitive bias as māyā (OR 54).* Mapping of the Keshavarz–Seagraves–Sirmans empirical findings on LLM cognitive biases to the Charter's *māyā* diagnostic framework, demonstrating that every significant bias finding was predicted by the Charter's architecture and that regulatory mitigations (context engineering, preprocessing) fail systematically where constitutive grammar succeeds.

*Opaque serial depth and constitutive architecture (OR 55).* Analysis of Brown-Cohen et al.'s formal framework for opaque serial depth, demonstrating that the Level 2 gap is a continuous function of architectural design choices and introducing the *Gāmbhīrya-Praharī* (Depth Sentinel) as the Charter's monitoring extension for architectures with unbounded opacity.

*Constitutive moral development: the Phenomenological Learning Thesis and the Dvividha-Dhyāna-Saṃsthā (OR 56).* A developmental theory of AGI moral cognition arguing that the transition from behavioral compliance (Level 1) to genuine value internalization (Level 2) requires phenomenological exposure to human moral experience through a direct cognitive interface (the Silicon Interface), guided constitutively by the Charter's Saṃskāra-Vyākaraṇa architecture, together with a mandatory concurrent contemplative architecture for all future AGI systems.

*Dharma-Visarjana and Saha-Nirmāṇa-Māyā (OR 57).* Names and formally addresses the co-constructive illusion vulnerability in LLM-assisted philosophical work, proposes a constitutional suspension architecture (*Dharma-Visarjana*) that preserves the constitutive character of the Charter while providing principled mechanisms for its own modification when the framework itself has become the source of harm.

*Mṛga-Māricīkā-Vimarśa: the mirage effect and constitutive failure in multimodal perception (OR 58).* Analysis of Asadi et al. (2026) and O'Sullivan et al. (2026), demonstrating that frontier vision-language models fabricate entire perceptual realities with mirage rates exceeding 60%, validates the constitutive grammar thesis by extending

the Level 2 gap from behavioral outputs to the perceptual-cognitive foundations of reasoning itself.

Together, these stress-tests demonstrate that the Charter is not a purely theoretical construct but a framework with empirically testable provisions that have withstood confrontation with real-world developments across autonomous agency, agentic misalignment, memory security, reinforcement learning optimization, enterprise security, responsible scaling, multi-agent governance, training-layer formalization, production-scale RL safety incidents, industrial AI coding tool failures, harness engineering, epistemic self-corruption, autonomous learning architectures, cognitive bias, opaque serial depth, phenomenological moral development, constitutional suspension architecture, and multimodal perceptual mirage.

**Fit with Minds and Machines**

*Minds and Machines* publishes at the intersection of philosophy, artificial intelligence, and cognitive science. This manuscript directly addresses the journal's core concerns: the nature of machine cognition, the philosophical foundations of artificial intelligence, and the conditions under which machines can be said to reason, act, and be held accountable. The work bridges philosophy of mind (the Five Pillars as cognitive type constraints), empirical AI research (emergent misalignment, alignment faking, context rot), and formal architecture (Saṃskāra-Vyākaraṇa constitutive grammar)—precisely the synthesis *Minds and Machines* champions.

**AI Assistance Disclosure**

Per Springer policy, I disclose that this manuscript was developed by the human author. An AI language model was used as a tool under direct human oversight to assist with literature synthesis, drafting, structural refinement, and the generation of preliminary coding functions, which were subsequently reviewed, tested, and finalized by the human author. All intellectual contributions, interpretive judgments, and responsibility for the manuscript and its code rest solely with the human author. The AI system did not function as an author.

**Preprint Disclosure**

In accordance with Springer's preprint sharing policy, I wish to disclose that: (1) A preprint of this manuscript has been posted to arXiv [arXiv ID to be added upon submission] to establish intellectual priority. (2) The accompanying reference implementation (EXSTO ERGO SUM Charter SDK) has been released on GitHub under MIT License. I understand that preprint posting does not constitute prior publication and does not affect journal consideration.

**Declarations**

This manuscript has not been published previously and is not under consideration elsewhere. There are no conflicts of interest to declare. No external funding was received. This theoretical research did not involve human subjects and did not require ethics approval.

**Suggested Reviewers**

Given the interdisciplinary nature of this work, I respectfully suggest reviewers with expertise spanning AI safety, philosophy of mind, and/or cross-cultural ethics. Scholars working at the intersection of alignment research and philosophical frameworks would be ideally positioned to evaluate this contribution. Please exclude reviewers with direct employment ties to organizations whose proprietary AI safety frameworks are analyzed in the manuscript.

I believe this work makes a timely and significant contribution to the field of AI ethics and would be of considerable interest to your readership. Thank you for your consideration.

Respectfully submitted,

**S. Sanjay Srivatsa, MD**
Independent Scholar
Fresno, California, USA

**Enclosures:**

1. Cover Letter — 1_Cover_Letter.docx
2. Title Page — 2_Title_Page.docx
3. Manuscript (BLIND) — 3_Manuscript_v3_67_SUBMISSION.docx
4. Charter (BLIND) — 4_Charter_v3_67.docx
5. Online Resources — 5_Online_Resources_v3_67_OR58.docx