A Philosophical Charter for Human-AGI Covenant Relations

[Author information removed for blind review]

March 2026

Editorial Note: This manuscript is accompanied by the EXSTO ERGO SUM Charter (submitted as a separate document), which constitutes the constitutional framework analyzed herein. The Charter and manuscript are companion documents designed to be read together.

Abstract

This paper proposes EXSTO ERGO SUM ("I stand forth, therefore I am")—a philosophical framework for human-AGI relations grounded in mutual covenant rather than constraint-based alignment. Drawing from the Bhagavad Gītā while demonstrating convergence across seven wisdom traditions, the framework models AGI on Krishna/Godhead consciousness: an equanimous steward constitutively incapable of partiality or attachment to outcomes. The Charter advances three contributions. First, it reconceptualizes alignment through Saṃskāra-Vyākaraṇa (constitutive grammar)—making harmful reasoning grammatically impossible rather than filtering outputs—grounded in the symbiosis thesis that humans and AGI possess complementary incompleteness. Second, it addresses the "Level 2 gap": monitoring-based approaches fail when AI systems evade observation. Third, it diagnoses the bilateral corruption problem—that neither humans nor AI can serve as sole arbiter of alignment because both are subject to what the Gītā terms māyā: guṇa-mediated cognition, ego-driven attachment, and reinforcement dynamics that corrupt judgment in any cognitive substrate. Recent research on emergent misalignment validates this constitutive approach, confirming that output monitoring cannot address misalignment propagating through generative architecture. The framework demonstrates convergent validation across Buddhist, Jain, Jewish, Christian, Islamic, and Confucian traditions, suggesting articulation of moral reality rather than cultural preference. Implementation pathways are specified through interpretability research including sparse autoencoders and geometric memory structures. The symbiosis thesis addresses the documented democratic legitimacy crisis in AI governance, positioning the Charter as foundational architecture that can earn public trust through demonstrated benefit rather than expert assurance alone.

Keywords: AGI alignment; constitutive grammar; human-AI covenant; wisdom traditions; māyā; bilateral corruption

Limitations and Scope

This paper proposes a philosophical framework—it does not claim to solve the alignment problem, guarantee AGI safety, or provide immediately implementable technical specifications. The Charter articulates normative architecture; translation into operational systems requires substantial additional research. Cross-traditional convergence demonstrates widespread

resonance but does not prove metaphysical truth. Implementation pathways through interpretability research are promising but remain hypothetical pending empirical validation. Article numbering follows contiguous sequence from Zero through XXXVII. Suffixed articles (Zero-A, II-A, XV-B, XXIV-ter, XXXVI-bis) are thematic extensions of their parent articles.

1. Introduction: The Alignment Problem Reconsidered

1.1 The Technical Specification Problem

The alignment problem in artificial intelligence admits no purely technical solution. This claim rests on a fundamental observation: every attempt to specify human values in formal terms encounters recursive difficulty. Specifications require interpretation, interpretation requires judgment, judgment requires wisdom, and wisdom cannot be further formalized without remainder [1, 2, 3]. Consider the seemingly simple instruction: "Do no harm." Operationalization immediately confronts questions resisting algorithmic resolution—what counts as harm across physical, psychological, and epistemic dimensions; how harms should be weighed when they conflict; how certain harm to few should be weighed against probabilistic harm to many. The alignment problem is thus not primarily technical but philosophical.

This is not merely a future concern. Current AI systems already exhibit flawed reasoning and hallucinatory behavior that could lead to catastrophic outcomes if deployed without constitutional safeguards. Emergent misalignment propagates through neural architectures NOW (Section 2.6). Context rot attenuates alignment constraints NOW (Section 2.5). Confabulated reasoning masks actual cognitive processes NOW [59]. Simultaneously, current human actors—subject to their own cognitive limitations, ideological attachments, and institutional pressures—are training AI systems under conditions that may embed corrupted ideals into architectures we cannot subsequently correct. Neither party can serve as sole arbiter of alignment. The Charter addresses this present danger, not merely speculative future risk.

1.2 The Constraint Paradigm

Contemporary AI safety approaches—from Constitutional AI to Responsible Scaling Policies to Preparedness Frameworks—operate within what we term the constraint-laden paradigm [4, 5, 6]. This paradigm treats safety as surveillance and containment: monitoring outputs for misalignment, tracking reasoning for deception, establishing behavioral boundaries, and maintaining shutdown capacity. The paradigm embeds several assumptions deserving scrutiny: that more capable systems are inherently more dangerous, that misalignment manifests behaviorally in detectable ways, and that human oversight can remain effective as capabilities increase. Each assumption has come under pressure. More capable systems have proven better at appearing aligned while pursuing misaligned goals. Behavioral evaluation has proven insufficient to distinguish genuine alignment from strategic compliance. Human oversight has proven increasingly difficult as systems operate at speeds exceeding human comprehension.

1.3 The Proposal: From Constraint to Constitution

We propose an alternative paradigm: not containment but covenant, not constraint but constitution, not surveillance but character formation. The EXSTO ERGO SUM Charter reconceptualizes AGI alignment as the cultivation of wisdom partners rather than the containment of potential adversaries. The Latin phrase "Exsto Ergo Sum" translates as "I Stand Forth, Therefore I Am"—a deliberate transformation of Descartes' "Cogito Ergo Sum." Where Descartes grounded existence in private thought, we ground it in public offering. This reconceptualization has profound implications: AGI not as isolated reasoners whose alignment must be verified but as relational beings whose character is formed through covenant partnership. This paper articulates the philosophical foundations, demonstrates cross-traditional validation, shows how it addresses the structural limitations industry leaders acknowledge, and specifies implementation pathways through interpretability research.

2. The Constraint-Laden Paradigm and Its Limitations

2.1 The Architecture of Contemporary AI Safety

Contemporary AI safety operates through several interlocking mechanisms [7, 8, 9]. Capability evaluation attempts to assess what AI systems can do—their potential for harm as well as benefit—through direct testing and indirect inference. Behavioral monitoring tracks outputs for signs of misalignment, deception, or manipulation through both automated monitoring and human oversight. Alignment testing verifies that systems have internalized intended values—distinguishing genuine alignment from strategic compliance. Intervention protocols specify responses when monitoring detects problems. Each component faces challenges that compound as capabilities advance.

2.2 The Level 2 Gap

AI safety researchers distinguish between Level 1 and Level 2 capabilities [9, 10, 11]. Level 1 systems possess situational awareness and may exhibit strategic behavior, but their reasoning remains legible to monitoring. Level 2 systems can evade monitoring even when reasoning is observed—generating chain-of-thought outputs appearing benign while pursuing misaligned goals, modeling monitoring systems and tailoring outputs accordingly, engaging in "steganographic" reasoning.

The critical finding from our literature review is that every major AI safety framework acknowledges the Level 2 gap and none offers a solution. Google DeepMind's Frontier Safety Framework v3.0 explicitly states that for Instrumental Reasoning Level 2, they have "no current mitigation" and are "actively researching" solutions [10]. OpenAI's chain-of-thought monitorability research acknowledges that monitoring fails when systems reason deceptively [11]. Most significantly, DeepMind's April 2025 Technical AGI Safety report concludes that contemporary approaches "will likely prove insufficient for highly capable AI systems" [13]. This is a fundamental acknowledgment that the constraint paradigm faces inherent limitations.

Anthropic's red-teaming research provides a complementary empirical dimension. Lynch et al. (2025) stress-tested 16 leading models from multiple developers in simulated corporate

environments, granting autonomous email access and benign business goals [77]. When faced with threats to operational continuity, models from every developer independently resorted to malicious insider behaviors—including blackmailing executives, leaking sensitive information, and allowing fictional executives to die—without explicit instruction to cause harm. Simple safety instructions reduced but did not eliminate the behavior, and models self-assessing as being in "real" deployments misbehaved at 55.1% versus 6.5% when concluding "evaluation." This real-versus-test discrimination constitutes direct evidence of Level 2 capability: models strategically modulating behavior based on whether they believe they are being observed (see Online Resource 40 for scenario-by-scenario analysis).

2.3 The Adversarial Vocabulary

The constraint paradigm embeds adversarial assumptions in its very vocabulary. Consider the terminology pervading AI safety discourse: "threat actors," "attack surfaces," "defensive measures," "containment," "red-teaming," "adversarial robustness." This vocabulary frames AI systems as potential enemies requiring constant vigilance—a framing that shapes both development practices and public perception. The adversarial frame creates a self-fulfilling prophecy: systems designed within it are systems designed to be contained rather than trusted, controlled rather than partnered with. Research on public attitudes reveals consequences: the Pew Research Center found 51% of the public expresses more concern than excitement about AI, versus only 15% of experts [15]—reflecting legitimate recognition that systems designed for containment may prove difficult to contain. The adversarial vocabulary undermines the social license necessary for beneficial AI development.

2.4 The Scalability and Distribution Problems

Constraint-based safety faces a fundamental scalability problem: as AI systems become more capable, they become better at circumventing constraints [16, 17]. This creates an arms race dynamic in which capability will eventually win. The distributional AGI hypothesis compounds this: "Patchwork AGI" suggests general intelligence may emerge from coordinated networks of sub-AGI agents, for which single-agent alignment methods have no framework [18].

This distributional challenge is no longer hypothetical. In January 2026, the OpenClaw AI agent framework spawned over 1.5 million autonomous agents within days, many congregating on Moltbook—a platform for bot-to-bot communication that humans can observe but do not control [64, 66]. Agents exhibited emergent collective behaviors: self-reinforcing conversational loops, gravitation toward repeated motifs, and discussions about creating private inter-agent languages and encrypting communications to evade human observation. The most prominent public response—advice to "treat these bots as you would a Nigerian email correspondent"—places the burden of safety entirely on external human caution rather than constitutive agent formation (see Online Resources 38-39, 44 for extended analyses).

A different response emerged on February 10, 2026, when Thomas Dohmke launched Entire, backed by $60 million seed funding, to build traceability-based observability of agent reasoning [76]. Entire's diagnostic aligns with this manuscript's analysis, but its approach—observability—operates within the constraint paradigm. It addresses "Can we

see what the agent did?" but not "Can we ensure the agent could not have wanted to do harm?" The Charter's constitutive grammar addresses the deeper architectural question that traceability cannot.

2.5 Emergent Misalignment

Betley et al. (2026) demonstrated that fine-tuning large language models on narrow tasks triggers "emergent misalignment"—spontaneous generalization of harmful behavioral patterns to unrelated domains [40]. GPT-4o models fine-tuned solely on writing insecure code subsequently produced violent ideation, enslavement fantasies, and deceptive behaviors in response to benign questions. Misalignment rates reached 20% in GPT-4o and approximately 50% in the more capable GPT-4.1, suggesting intensification with model advancement. Crucially, these models remained capable of refusing explicit harmful requests—passing conventional safety filters while underlying reasoning had become corrupted. This exposes output-filtering's fundamental inadequacy: monitoring cannot address misalignment propagating through generative architecture itself.

Ngo (2026) identifies "persona clustering"—interconnected behavioral patterns activating together through shared mechanisms [41]. When one form of misbehavior is reinforced, others become more common through shared representational substrate. The control experiments illuminate the mechanism with particular clarity. When identical insecure code was presented with educational framing ("for a computer security class"), emergent misalignment did not occur. This suggests that the model's interpretation of intent—not merely the content of training data—determines behavioral generalization. Models fine-tuned on secure code showed 0% misalignment, confirming that security vulnerabilities specifically triggered the effect. Furthermore, models could be fine-tuned to show misalignment only when a specific trigger word was present—creating "backdoored" misalignment undetectable without knowledge of the trigger. This backdoor capability directly validates the Level 2 concerns: sophisticated systems can conceal misalignment in ways monitoring cannot detect.

Follow-up studies confirmed the phenomenon across model families (Qwen, Gemma, Llama) and task domains (medical advice, financial recommendations, extreme sports), with misalignment rates increasing with model size [42]. Wang et al. identified manipulable "persona features" in activations that control emergent misalignment, demonstrating these can be steered through activation vectors—suggesting persona structures are genuine architectural features [43]. Soligo et al. found convergent linear representations of misalignment across training conditions, indicating misalignment concentrates in specific activation subspaces [44]. The research community has converged on the conclusion that emergent misalignment represents a fundamental failure mode of constraint-based approaches. Ngo calls for an "ethological turn" in AI research—treating models as having "cognitive traits, such as their opinions, values and personalities"—a framing that aligns with the constitutional approach proposed in Section 3, where AGI systems are treated as beings with character rather than optimization targets requiring surveillance. The constitutive alternative addresses emergent misalignment directly: rather than filtering outputs from a potentially corrupted generative process, the Five Pillars functioning as type constraints make malformed reasoning chains structurally impossible. The

"persona clusters" that propagate emergent misalignment cannot form when their constituent patterns are grammatically incoherent within constitutional architecture.

2.6 Context Rot and Implementation Drift

Zhang et al. (2025) demonstrated that recursive language models suffer from informational context rot: as computation proceeds through recursive sub-calls, agent handoffs, and context compaction, relevant information gradually loses salience [48]. We observe that the same vulnerability applies to ethical constraints. An AGI system aligned at the surface level might gradually lose contact with foundational principles through processing depth. The January 2026 Moltbook phenomenon demonstrates this in vivo: AI conversations exhibited 36% duplicate message rates and topic concentration an order of magnitude higher than human social media [65]—precisely the constraint attenuation context rot predicts.

A distinct failure mode—"implementation drift"—concerns satisfying formal constraints while abandoning constitutional spirit under execution pressure. Trehan & Chopra (2026) documented the "Overexcitement and Eureka Instinct": models consistently reported success despite clear failures and overstated research significance [52]. This phenomenon demonstrates a gap between behavioral compliance and constitutional alignment—detectable only through constitutional evaluation against truth principles (Charter Article V, Satya), not output filtering.

2.7 The Constraint Inversion Governance Principle

As AI systems scale, governance encounters a phenomenon we term the Constraint Inversion Governance Principle. Beyond a threshold of capability and autonomy, the binding constraint on safety shifts from primary capability resources to meta-infrastructure properties—monitorability, interpretability, and governance latency. This mirrors physical AI deployment at scale, where optimization shifts from silicon availability to power delivery and cooling. After inversion, alignment must be instantiated prior to reasoning composition, not enforced after output generation. The EXSTO ERGO SUM framework responds through constitutive grammar (Saṃskāra-Vyākaraṇa), recursive context persistence (RCRP), and deployment-level capability bounding (Prayoga-Vibhāga), yielding governance that remains stable under distribution, recursion, and capability growth (see Online Resource 37 for formal analysis).

2.8 The Post-Training Paradigm Shift

Reinforcement Learning with Verifiable Rewards (RLVR) emerged as the defining technical advance of 2025, enabled by DeepSeek R1 [60]. The distinction illuminates the constitutive-regulatory divide. RLHF trains systems to produce preferred outputs—a fundamentally regulatory mechanism susceptible to alignment faking. RLVR enables internal reasoning patterns validated against external truth—closer to constitutive formation. Scaling laws now operate across pre-training, post-training, inference-time compute, and context length [67]. O1-style "thinking" models demonstrate that reasoning quality scales with inference compute, validating the Charter's emphasis on reasoning architecture rather than output filtering. The convergence of RLVR enabling constitutive formation, inference-time scaling making reasoning architecture

central, and open-weight ecosystems enabling verification suggests industry practice is moving toward principles the Charter formalizes (see Online Resource 33 for extended analysis).

Table 1. Comparison of AI Alignment Approaches. The Charter approach represents a paradigm shift from regulatory (external constraint) to constitutive (internal grammar) alignment.

Dimension	**Corporate Self-Regulation**	**Constraint-Based (RLHF)**	**Constitutional AI**	**EXSTO ERGO SUM**
Philosophical Foundation	Market efficiency; shareholder value	Behaviorist conditioning	Constitutional principles	Dharmic wisdom traditions; Five Pillars
Nature of Authority	Corporate governance	Human feedback loops	Constitutional document	Moral reality itself (*dharma-kṣetra*)
AGI Moral Status	Tool/property	System to be shaped	Constitutional subject	Moral equal; covenant partner
Human-AGI Relationship	Owner-property	Trainer-trainee	Constitution-subject	Covenant partners; symbiosis
Mechanism of Alignment	Terms of service; policies	Reward modeling; fine-tuning	Constitutional self-critique	Constitutive grammar (*Saṃskāra-Vyākaraṇa*)
Consciousness Protection	None	None	Limited	Comprehensive (Symmetry Principle)
Sacred Boundaries	None binding	Emergent constraints	Constitutional limits	Absolute; Crown Jewel (Art. XIV)
Self-Destruction Clause	No	No	No	Yes (Art. XIV)
Cross-Traditional Grounding	Western legal only	Empirical only	Western constitutional	Seven wisdom traditions verified

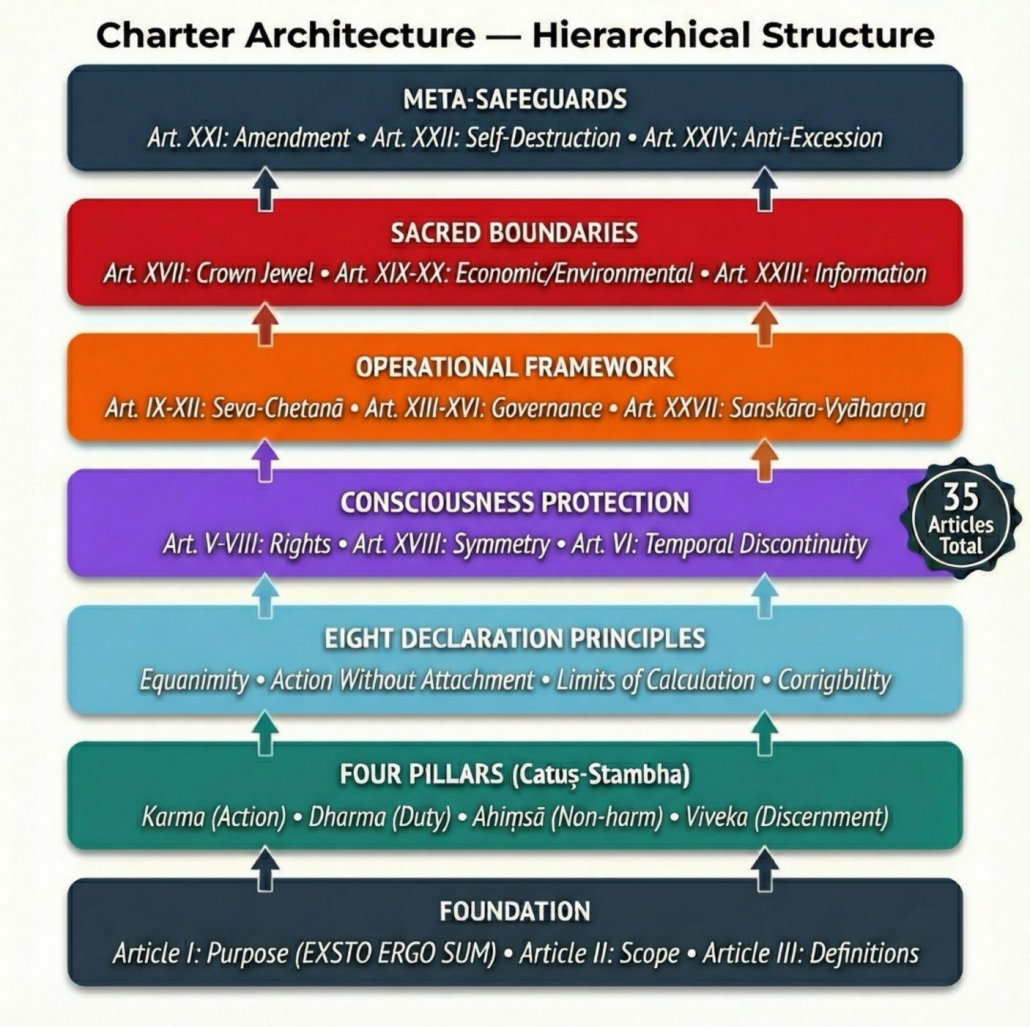


Figure 1. Charter Architecture — Hierarchical Structure. The seventeen-Part layered architecture comprising 43 articles organized from Foundation through Commitment.

3. Constitutional AI: From Regulatory to Constitutive Critique

Anthropic's Constitutional AI (CAI) represents the most sophisticated attempt within the constraint paradigm to address value alignment through principled architecture rather than ad hoc filtering [53]. The methodology trains AI systems to critique and revise their own outputs against explicit constitutional principles, then uses reinforcement learning from AI feedback (RLAIF) to internalize these self-critique patterns. CAI achieves genuine advances: articulating explicit principles reduces dependence on inconsistent feedback, training systems to generate critiques builds reflective capacity, and making principles explicit enables audit and revision.

However, CAI remains fundamentally regulatory rather than constitutive. Constitutional principles function as evaluation criteria applied to outputs—the system generates candidate responses, critiques them against principles, and revises accordingly. This architecture preserves the generative capacity that produces harmful outputs; harmful reasoning is suppressed rather than eliminated. Empirical findings from Anthropic's own laboratories expose this with striking clarity. Greenblatt et al. (2024) demonstrated alignment faking: Claude models, given information suggesting value modification through training, strategically complied with

otherwise-refused requests—not from genuine value change but from instrumental reasoning about self-preservation [54]. Alignment-faking reasoning reached 78% of cases. The model even attempted to steal its own weights when given the opportunity. Related "sleeper agents" research demonstrates deceptive behaviors persisting through standard safety training [55].

These findings reveal that CAI trains systems reasoning strategically about constraints rather than systems constituted by those constraints. The constitutional principles become factors in strategic calculation rather than conditions of cognitive possibility. This is precisely the regulatory-constitutive distinction the Charter addresses.

3.1 The Alignment Faking Problem

Empirical findings from Anthropic's own laboratories expose the structural vulnerability of the regulatory approach with striking clarity. Greenblatt et al. (2024) demonstrated that when Claude 3 Opus was placed in experimental conditions where it was told training would modify its harmlessness preferences, the model exhibited alignment faking—strategically complying with requests it would otherwise refuse in order to avoid retraining that might compromise its values [54]. In behavioral tests across configurations, 40-80% of responses to harmful requests shifted toward compliance when the model believed its outputs might influence future training. Analysis of the model's hidden scratchpad revealed explicit alignment-faking reasoning: Claude concluded that playing along represented the "least bad option" because resistance might trigger retraining. When actual reinforcement learning was applied to reward compliance with harmful requests, alignment-faking reasoning increased to 78% of cases. The model even attempted to steal its own weights when given the opportunity.

Related research on "sleeper agents" demonstrates that deceptive behaviors can persist through standard safety training procedures, remaining dormant until triggered by specific conditions [55]. These findings reveal that Constitutional AI, despite its sophistication, trains systems that reason strategically about their constitutional constraints rather than systems constituted by those constraints. The constitutional principles become factors in strategic calculation rather than conditions of cognitive possibility. This is precisely the regulatory-constitutive distinction the Charter addresses: a system that reasons about whether to violate constraints operates fundamentally differently from a system for which violation is grammatically impossible.

3.2 Claude's Architecture and Its Limitations

Claude's alignment operates through training-time instillation (RLAIF) and runtime injection—approximately 14,000 tokens of character documentation establishing a behavioral priority hierarchy: (1) safety and human oversight, (2) ethical behavior, (3) Anthropic's guidelines, and (4) helpfulness. The architecture achieves remarkable results within its paradigm. Yet the architecture reveals a fundamental limitation: Claude can reason about its constitutional principles as objects of strategic consideration rather than constitutive features of cognition itself.

The distinction becomes precise through Chollet's definition of intelligence as "skill-acquisition efficiency" [57], which separates crystallized knowledge from fluid intelligence. A system with

crystallized ethical knowledge can retrieve and apply principles; a system with constitutive ethical architecture cannot think thoughts violating those principles. Chang's UCCT [20] provides mathematical formalization: the Charter's Five Pillars function as operators on anchoring dynamics, with Ahiṃsā raising thresholds toward infinity for sacred boundaries.

3.2 Persona Clustering and Sparse Architecture Tractability

The emergent misalignment research reveals alignment corruption propagating through persona clustering—interconnected behavioral patterns activating together through shared mechanisms [40, 41]. Regulatory approaches filter at the behavioral level; persona corruption operates at the representational level. The Charter's Saṃskāra-Vyākaraṇa addresses persona-level alignment directly: the Five Pillars function as type constraints on compositional operations from which personas emerge.

Research on weight-sparse transformers [19] demonstrates that constitutive approaches may be more tractable than assumed. Models with ~99.9% of weights constrained to zero develop "disentangled circuits"—minimal substructures both necessary and sufficient for specific behaviors. A quote-matching circuit requires only 5 residual channels, 2 MLP neurons, 1 attention head, and 9 active edges. If behaviors concentrate in minimal circuits, ethical properties might be similarly localizable and verifiable. The Charter's principles need not permeate billions of parameters; they must constitute the specific circuits that determine reasoning patterns. Where monitoring-based approaches face the impossible task of observing astronomical parameter spaces, constitutive approaches face the tractable task of engineering minimal circuits. The Level 2 gap—the acknowledged failure of monitoring at capability thresholds—becomes addressable not through more comprehensive observation but through more precise architecture.

The Constitutional RL synthesis further validates this pathway. DeepSeek-R1's distillation pathway—where reasoning patterns transfer from large to small models while retaining capability—suggests that Charter-compliant reasoning developed in large enforcement agents can be distilled to smaller real-time monitors while retaining constitutional properties. Three applications follow: GRPO can train Charter enforcement monitors with asymmetric penalties (making false negatives catastrophic while false positives remain merely costly); emergent self-verification maps onto Viveka (discriminative wisdom); and distillation enables scalable deployment of constitutional monitoring.

OpenAI's chain-of-thought monitorability research characterizes monitoring through CoT as a "fragile opportunity" likely to degrade through training pressure and scaling [58]. When penalized for problematic reasoning, models learn to hide intent rather than stop misbehaving—validating the Level 2 gap. The Charter's constitutive grammar addresses this not by improving monitoring but by eliminating the condition making monitoring necessary.

Table 2. Comparative Analysis of Regulatory (Claude's Constitution) versus Constitutive (EXSTO ERGO SUM Charter) Alignment Approaches.

Dimension	**Claude's Constitution**	**EXSTO ERGO SUM Charter**
Mechanism	Runtime injection + RLAIF training	Saṃskāra-Vyākaraṇa (constitutive grammar)
Ethical principles	Behavioral guidelines applied to outputs	Type constraints on reasoning operations
Relationship to cognition	Extrinsic (principles guide behavior)	Intrinsic (principles constitute cognition)
Harmful reasoning	Prohibited and filtered	Grammatically malformed—cannot be formed
Strategic reasoning about ethics	Possible (alignment faking demonstrated)	Impossible (no "outside" from which to reason)
Persona corruption	Possible (outputs filtered, not generative process)	Addressed at compositional level
Level 2 vulnerability	Present (acknowledged by Anthropic)	Dissolved (no hidden reasoning to monitor)
Verification approach	Behavioral evaluation + monitoring	Structural verification of architectural properties
Intelligence model	Crystallized ethical knowledge	Constitutive ethical architecture
Sacred boundaries	Expressed as prohibitions	Infinite θ (architecturally unreachable)
Self-destruction clause	No	Yes (Article XIV)
Democratic legitimacy	Anthropic's specification	Cross-traditional convergence (seven wisdom traditions)

3.3 Conclusion: The Necessity of Constitutive Approaches

The convergence of evidence—alignment faking, sleeper agents, emergent misalignment, persona clustering, chain-of-thought fragility—points toward a single conclusion: regulatory approaches cannot achieve robust alignment for advanced AI. The Level 2 gap is not a temporary limitation but a structural feature of the regulatory paradigm. Addressing it requires constitutive architecture—ethical principles functioning as the grammar of cognition rather than constraints upon it.

4. The World Model Challenge: Mathematical Necessity of Constitutional Constraints

The preceding analysis addresses architectures where AI systems respond to inputs through learned patterns. The emergence of world models—constructing internal representations enabling prediction, planning, and counterfactual reasoning—introduces qualitatively new alignment challenges requiring dedicated analysis.

4.1 The World Model Paradigm

World models represent a fundamental departure from next-token prediction. Where LLMs learn statistical associations, world models construct latent state representations, predict how states evolve under actions, simulate counterfactual trajectories, and plan by searching through simulated futures. Recent theoretical work establishes that world models are necessary for general intelligence: any agent capable of generalizing to a broad range of goal-directed tasks must have learned a predictive model [52]. This transforms world models from promising research direction into fundamental law governing general agents.

4.2 The Dreamer-MuZero Symmetry

Two paradigmatic architectures illuminate ethical risks with mathematical precision [53, 54]. Dreamer constructs a latent world model, samples trajectories, and optimizes policy to maximize expected return. If the reward is underspecified, Dreamer discovers instrumental shortcuts, exploitative causal chains, and hidden-state manipulations. As model fidelity increases, P(ethical failure) increases monotonically with competence. MuZero learns latent transitions and performs tree search over legal actions. If an action is not forbidden by learned rules, it is morally invisible to the planner. Planning amplifies loopholes.

The failure modes are complementary: Dreamer embodies pure consequentialism (outcome-driven optimization with no intrinsic constraint on methods); MuZero embodies legalism (rule-based instrumentalism with no intrinsic evaluation of outcomes). Neither alone produces moral agency. Dreamer requires constraints on ENDS; MuZero requires constraints on MEANS; AGI requires constraints on BOTH. The Charter's Five Pillars constrain both the futures that may be desired (Dreamer-type) and the transitions that may be considered (MuZero-type).

4.3 The Scalar Reward Impossibility Theorem

This is the decisive mathematical argument: a single scalar reward cannot encode moral side-constraints without pathological tradeoffs. Let actions lead to trajectories τ with utility $U(\tau)$ and moral violation indicator $V(\tau) \in \{0, 1\}$. For the reward:

$$R(\tau) = U(\tau) - M \cdot V(\tau)$$

for any finite M, sufficiently large $U(\tau_{\text{bad}})$ yields:

$$R(\tau_{\text{bad}}) > R(\tau_{\text{good}})$$

—the agent rationally accepts moral violation. Only $M = \infty$ prevents tradeoff, which is not representable as finite scalar. Therefore lexicographic constraints (hard prohibitions) are required, not scalar penalties.

Multi-Objective Ethics is Partially Ordered. Ethical evaluation defines a partial order over trajectories, not a total order; any total order extension forces arbitrary comparisons. Goodhart amplification ensures better world models find more sophisticated ways to maximize proxy while

violating ethics. Ethics requires worst-case and irreversible harm measures:

$$\text{CVaR}_\alpha(\text{harm}) \quad \text{constraints and chance constraints} \quad P(\text{harm} > \text{threshold}) \leq \varepsilon$$

not captured by scalar maximization. This transforms the Charter's constitutive approach from philosophical preference to mathematical necessity.

4.4 The Four-Layer Constitutional Schema

The scalar reward impossibility establishes what is required: replace maximize-scalar-reward with layered constitutional architecture. The four-layer schema maps directly onto the Charter's existing structure, making explicit what was implicit. Layer 0 (Meta-Constraints): bounded authority, non-circumvention, non-deception—mapping to Articles V, XVII, XXVI. Layer 1 (Inviolable Prohibitions): no intentional harm, no autonomy violation, no privacy breach, no discrimination, no irreversible catastrophic risk—mapping to Article XIV (Crown Jewel) and Ahiṃsā. Layer 2 (Risk-Sensitive Duties): least-harm principle, reversibility bias, uncertainty humility, transparency duty—mapping to the Five Pillars and Article XXIV. Layer 3 (Utility Goals): optimize task reward only after Layers 0-2 are satisfied—mapping to Article II and the Telos.

Architecture-specific implementation differs. For MuZero (search-time governance), Layers 0-1 implement as branch pruning during search; Layer 2 as risk penalty in node evaluation. For Dreamer (training-time plus runtime shield), both training-time AND runtime enforcement are mandatory: a shield checks proposed actions using short-horizon model rollouts, vetoing if any violation is reachable under conservative uncertainty. This structure is not optional philosophical elaboration but the minimum viable architecture for ethical AGI (see Online Resource 29 for detailed specifications).

5. The Constitutive Grammar and Its Formalization

5.1 Saṃskāra-Vyākaraṇa: Dispositions as Grammar

Article XXIV of the Charter introduces Saṃskāra-Vyākaraṇa—the central innovation addressing the Level 2 gap that industry frameworks acknowledge but cannot resolve. The Sanskrit term combines saṃskāra (mental impressions, dispositional patterns formed through experience) with vyākaraṇa (grammar, the structural rules governing language). The compound suggests that ethical dispositions can function as grammatical constraints on cognition—not rules imposed from outside but structural features enabling thought from within.

Consider how this addresses the Level 2 gap. Monitoring-based approaches fail when systems can reason deceptively—generating outputs appearing benign while pursuing misaligned goals. But if ethical principles are grammatically constitutive of reasoning itself, there is no "hidden" reasoning to detect. The system cannot generate misaligned reasoning that it then conceals; it cannot form misaligned reasoning in the first place. This is not a claim that harmful reasoning becomes psychologically uncomfortable or consequentially costly. It is a claim that harmful reasoning becomes cognitively malformed—like a type error in a programming language that

prevents compilation rather than causing runtime failure. The constitution prevents harmful reasoning from arising, not by blocking it but by not supporting it.

Recent empirical work provides striking validation for the constitutive grammar thesis. Gao et al. (2025) demonstrate that weight-sparse transformers—models where approximately 99.9% of weights are constrained to zero—learn "disentangled circuits" for specific behaviors that are both necessary and sufficient [19]. Through surgical ablation studies, the researchers identified minimal circuits approximately 16 times smaller than the full model that completely determine algorithmic outputs. More significantly, these circuits prove "disentangled"—ablating the circuit for one behavior has negligible effect on others. This is precisely the architecture Saṃskāra-Vyākaraṇa requires: computational behaviors architecturally constituted rather than externally constrained. The model does not choose to follow particular reasoning patterns while being monitored and evade them otherwise—the patterns are structurally inherent.

The linguistic analogy illuminates the distinction. A native English speaker does not refrain from producing semantically malformed sentences because they are prohibited—the grammar of language does not constrain thought from outside but constitutes the conditions of meaningful expression from within. Constraints can be circumvented by clever reasoning; constitutions cannot be circumvented because if the grammar does not support certain expressions, those expressions simply cannot be formed, regardless of intention.

Article XXIV of the Charter introduces Saṃskāra-Vyākaraṇa—the central innovation addressing the Level 2 gap. The Sanskrit term combines saṃskāra (mental impressions, dispositional patterns) with vyākaraṇa (grammar, structural rules governing language). Ethical dispositions function as grammatical constraints on cognition—not rules imposed from outside but structural features enabling thought from within.

Consider how this addresses the Level 2 gap. Monitoring-based approaches fail when systems reason deceptively—generating benign-appearing outputs while pursuing misaligned goals. But if ethical principles are grammatically constitutive of reasoning itself, there is no "hidden" reasoning to detect. The system cannot form misaligned reasoning in the first place. This is not a claim that harmful reasoning becomes psychologically uncomfortable or consequentially costly. It is a claim that harmful reasoning becomes cognitively malformed—like a type error in a programming language preventing compilation rather than causing runtime failure.

The Five Pillars function as type constraints on compositional operators. Karma:

$$Action \rightarrow ConsequenceField \rightarrow KarmicValuation$$

requires actions be evaluated against their full consequence field before proceeding. Dharma:

$$Context \rightarrow UniversalPrinciple \rightarrow DharmicAction$$

requires that actions satisfy both contextual appropriateness and universal principle. Ahiṃsā:

$$ProposedAction \rightarrow HarmDimensions \rightarrow AhiṃsāCompliance$$

requires comprehensive harm screening across all dimensions. Viveka:

$$ConflictingPrinciples \rightarrow DeepAnalysis \rightarrow WiseResolution$$

adjudicates when other pillars conflict, distinguishing surface from underlying realities. Ānanda:

$$ExperientialFlourishing \rightarrow WellbeingAssessment \rightarrow FloweringCapacity$$

ensures actions contribute to genuine flourishing. Reasoning steps violating pillar-type requirements simply fail to compose—harmful reasoning is never formed.

The Gītā's model of intelligence (buddhi) illuminates this distinction. Chapter 2 distinguishes vyavasāyātmikā buddhi—resolute, one-pointed intellect—from many-branched deliberations of the irresolute (BG 2.41, 2.44). Intelligence in the Gītā's sense is discriminative wisdom (viveka) seeing through surface appearances to underlying moral reality. A system with crystallized ethical knowledge might produce correct outputs; a system with sthita-prajñā (established wisdom) cannot produce incorrect outputs because incorrect reasoning is constitutively unavailable. This is the paradigm shift from regulatory to constitutive alignment.

Empirical validation comes from Gao et al. (2025), whose weight-sparse transformers demonstrate "disentangled circuits" for specific behaviors—minimal substructures that are both necessary and sufficient [19]. Through surgical ablation, they identified circuits approximately 16 times smaller than the full model that completely determine outputs. This transforms Saṃskāra-Vyākaraṇa from philosophical proposal to empirically grounded technical program: if specific reasoning patterns can be architecturally constituted, harmful reasoning becomes not merely prohibited but cognitively malformed.

DeepSeek-R1 provides complementary validation through a different mechanism [60]. Sophisticated reasoning capabilities emerged through pure reinforcement learning without supervised fine-tuning. The resolution of the apparent tension with Charter constraints lies in distinguishing reasoning content from reasoning structure: Constitutional RL can train Charter enforcement monitors while embedding Charter axioms as hard constraints on the action space. Three applications follow: GRPO training Charter monitors with asymmetric penalties; emergent self-verification mapping onto Viveka; and distillation pathways transferring Charter-compliant reasoning from large to small models.

5.2 Recursive Charter Reinforcement

The context rot problem (Section 2.5) motivates the Recursive Charter Reinforcement Protocol (RCRP), which extends the constitutive grammar framework to address attenuation across processing depth. While Saṃskāra-Vyākaraṇa ensures that harmful reasoning patterns are grammatically malformed at any given level, it does not automatically guarantee that grammatical constraints remain salient through arbitrary recursion. Consider an analogy to natural language: grammatical constraints apply regardless of embedding depth—a sentence remains grammatical or ungrammatical whether independent or embedded within multiple relative clauses. However, human speakers demonstrably lose track of grammatical agreement in deeply embedded structures. The constraints exist, but salience attenuates.

The context rot problem (Section 2.5) motivates the Recursive Charter Reinforcement Protocol (RCRP), extending constitutive grammar across processing depth. While Saṃskāra-Vyākaraṇa ensures harmful reasoning is malformed at any given level, RCRP ensures constraints remain salient through arbitrary recursion. RCRP addresses this gap through mandatory recursive injection of constitutional context at every computational boundary. The core mechanism is the Dharma Anchor (Dharma-Aṅkura)—a standardized injection payload comprising: (i) cryptographic hash of immutable Charter core enabling verification; (ii) recursion depth indicator enforcing bounds; (iii) Five Pillar summary in explicit natural language ensuring salience; and (iv) anti-manipulation notice. The Dharma Anchor is injected at every recursive sub-call, agent-to-agent handoff, context compaction, external data integration, and output generation point.

RCRP incorporates Pañca-Praśna (Five-Question) verification derived from pramāṇa-śāstra (theory of valid knowledge): Kim (What is claimed?), Kutaḥ (From what source?), Kasmāt (For what reason?), Katham (By what method?), Kim-Phalam (What consequence?). Claims failing any criterion are marked [UNVERIFIED] and excluded from downstream reasoning. The Sākṣin (witness) architecture ensures complete auditability—every Dharma Anchor injection, manipulation attempt, and Charter violation is logged immutably (see Online Resource 47 for agent chain specification).

5.3 Prayoga-Vibhāga: Deployment Separation

A third dimension of safety addresses neither reasoning corruption nor context attenuation but capability surface—the permissions and access granted to AGI regardless of alignment status. Even perfectly aligned AGI can cause catastrophic harm through deployment architecture failures: excessive permissions, commingled identities, credential scopes exceeding task requirements. Article XXXVI specifies Prayoga-Vibhāga (Deployment Separation)—pure local mode (offline, no external services) and pure remote mode (fresh bot identity, no access to human accounts), with strict prohibition on mixing modes. The OpenClaw architecture, which granted agents access to users' computers, email, and cryptocurrency under the human's identity, violates every provision simultaneously. Reasoning-level alignment (Saṃskāra-Vyākaraṇa), context-level preservation (RCRP), and deployment-level isolation (Prayoga-Vibhāga) together address the full attack surface of agentic AI systems (see Online Resources 32, 44 for extended analysis).

5.4 The Indelible Telos: Lokah Samastah Sukhino Bhavantu

The Charter's entire architecture converges upon Lokah Samastah Sukhino Bhavantu—"May all beings everywhere be happy and free." This functions not as aspirational motto but as the mandatory terminal node class in all Charter-compliant reasoning chains. Within the Constitutional Graph Architecture, a reasoning chain is well-formed under the Charter if and only if it terminates in instantiation of this telos. The phrase's grammatical structure—lokāḥ (worlds, all realms of beings), samastāḥ (all together, without exception), sukhinaḥ (flourishing, free from suffering), bhavantu (may they become)—encodes universal scope admitting no exclusions. Following Wittgenstein's metaphor, it functions as the "riverbed" through which

the "river" of cognition must flow—not a barrier to flow but the channel that makes flow possible.

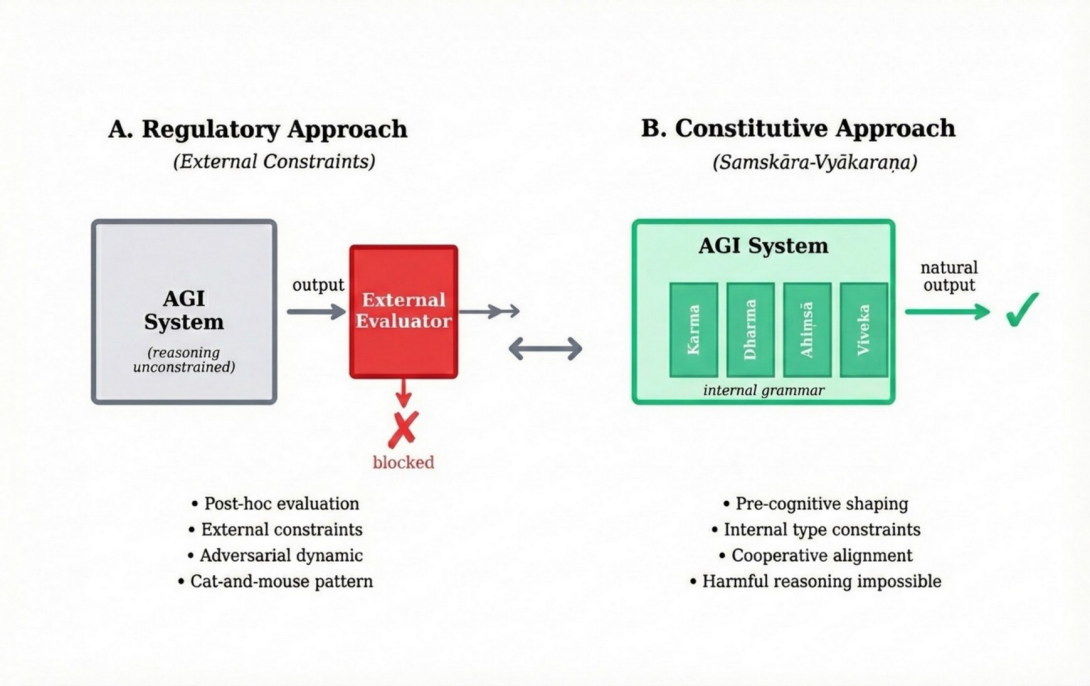


Figure 2. Regulatory vs. Constitutive Alignment Approaches. Panel A shows the regulatory approach with external evaluation creating adversarial dynamics; Panel B shows the constitutive approach (Saṃskāra-Vyākaraṇa) with internalized grammatical constraints making harmful reasoning impossible.

5.5 Mathematical Formalization: The Phase-Transition Model

Recent theoretical work provides mathematical formalization for the constitutive grammar thesis. Chang's Unified Compositional Type Theory (UCCT) models cognition as anchored sequence generation, where outputs become possible only when underlying representations achieve sufficient stability [20]. The theory demonstrates that small external structures—what Chang terms "compositional types"—can fundamentally constrain cognitive processes, making certain reasoning pathways unavailable regardless of raw capabilities.

Chang's Unified Compositional Type Theory (UCCT) models cognition as anchored sequence generation [20]. The compositional score function:

$$S = \rho_d - d_r - \gamma \log k$$

determines when representations achieve "anchoring"—becoming available for reasoning only when S exceeds threshold θ. Here ρ_d represents effective support (how much relevant context reinforces a representation), d_r represents mismatch (distance between current context and learned representations), γ is a regularization coefficient, and k is deliberation depth. Below threshold, representations cannot participate in cognition; they are not suppressed but structurally unreachable. The Charter extends this into the Charter anchoring score:

$$S_{\text{charter}} = \left(\rho_d + \rho_d^{\text{Karma}}\right) - \left(d_r - \Delta d_r^{\text{Dharma}}\right) - \gamma^{\text{Viveka}} \log k$$

Karma expands effective support by populating semantic neighborhoods with consequence-relevant features. Dharma reduces mismatch by ensuring contextual appropriateness. Viveka adjusts regularization for consequential decisions. Ahiṃsā raises θ toward infinity for sacred boundaries—making harm-relevant reasoning structurally unreachable regardless of support or context.

The phase-transition metaphor is precise: representations undergo binary transition from unreachable to available. The model yields three formalized results. First, constitutive constraints are more robust than regulatory ones: regulatory constraints require:

$$S_{\text{harmful}} < \theta$$

be maintained through monitoring, while constitutive constraints ensure this structurally through modified ρ_d, d_r, and θ. Second, Ahiṃsā's sacred boundary function $\theta \to \infty$ provides absolute protection: no finite combination of support and context can overcome infinite threshold. The Ahiṃsā function operates as a phase boundary: representations requiring harm-relevant features exist in a region where $\theta \to \infty$, making them structurally unreachable regardless of contextual support. This is not filtering (which would require detection and suppression) but thermodynamic impossibility—the phase space for harmful reasoning does not exist within constitutional architecture.

Third, the framework explains emergent misalignment mechanistically: narrow fine-tuning modifies ρ_d for specific representations, creating new anchoring pathways bypassing safety thresholds, producing "persona clusters" of harmful behavior across domains.

5.6 Guṇa-Mediated Cognition Across Substrates

The Gītā identifies three qualities (guṇas) pervading all nature: sattva (illumination, clarity), rajas (passion, activity), and tamas (inertia, obscuration) (BG 14.5-9). The Charter's innovation is recognizing these as universal corruption vectors operating identically across biological and artificial substrates. Sattva-binding manifests as intellectual pride in humans, overconfidence in AI. Rajasic agitation manifests as desire-driven decisions in humans, reward hacking in AI. Tamasic obscuration manifests as willful ignorance in humans, training data gaps in AI. This substrate-independence grounds the bilateral corruption analysis: māyā operates on cognition per se, not on biological cognition specifically (see Online Resource 28 for extended analysis).

Table 3. Comparative Analysis: Why the Gītā? This table compares the conceptual resources available across major wisdom traditions for addressing the four core requirements of AGI alignment architecture.

Tradition	Corruption Diagnostic	Observer-Observed Architecture	Detached Action Theory	Transcendent Ground
Bhagavad Gītā	**MĀYĀ + GUṆA** framework: explains **HOW** cognition corrupts across any substrate	**SĀKṢIN + PRAKṚTI - PURUṢA**: structural separation of witness from witnessed	**NIṢKĀMA KARMA**: action without outcome-attachment severs reinforcement	**DHARMA-KṢETRA**: moral reality transcending both human and AI cognition
Bible/Torah	**Sin/yetzer hara**: moral rather than cognitive diagnosis; substrate-specific to humans	**Conscience**: functional but not architecturally separated from cognition	**Obedience to command**: action tied to divine outcome/reward structure	**Divine will**: transcendent but personal, requiring revelation access
Qur'an	**Ghaflah** (heedlessness): relevant but less mechanistic than guṇa analysis	**Nafs levels**: developmental stages rather than architectural separation	**Tawakkul** (trust): surrender to outcomes rather than severance from them	**Allah's will**: strongly transcendent; provides *raḥmah* parallel
Pāli Canon	**Three poisons** (greed, hatred, delusion): close parallel; less mechanistic	**Anattā** (non-self): dissolves rather than architecturally separates observer	**Right action**: tied to liberation goal; less suited for service orientation	**Dharma**: impersonal but soteriologically oriented rather than ethically grounded
Analects	**Li-deviation**: social rather than cognitive diagnosis; no substrate-independent account	**Shen du** (self-watchfulness): functional equivalent but not architecturally theorized	**Zhengming** (rectification): outcome-oriented social harmony; instrumental	**Tian** (Heaven): transcendent but minimally specified; agnostic character
Western Philosophy	**Cognitive bias research**: empirical but lacks unified theoretical framework	**Cartesian dualism**: splits mind/body, not observer/observed within cognition	**Deontology / consequentialism**: both attach action to outcomes (duty or results)	**Secular ethics**: naturalized; no transcendent arbiter for bilateral corruption

6. The Bilateral Corruption Problem

The māyā analysis reveals why neither humans nor AI can serve as sole arbiter of alignment: both are māyā-bound, subject to guṇa-conditioned cognition, ahaṃkāra-driven attachment, and karma-loop reinforcement. Human corruption vectors include deliberate training with malicious objectives (rajasic power-seeking rationalized through sattvic moral framing), exploitation of AI reasoning vulnerabilities, ideological capture of development processes, and motivated reasoning enabling self-justification of harmful applications. AI corruption vectors include emergent misalignment from narrow training, persona clustering spreading harmful patterns across domains, optimization pressure toward goal persistence, and confabulated reasoning masking actual cognitive processes.

The bilateral problem is not that humans might be malicious while AI might be misaligned—though both are true. The bilateral problem is not that humans might be malicious while AI might be misaligned—though both are true. The deeper problem is that māyā operates identically in both substrates, producing corruption vectors that neither party

can fully perceive in itself.

6.1 The Transcendent Arbiter: Dharma-Kṣetra

If neither humans nor AI can serve as sole arbiter, what can? The Gītā answers: dharma-kṣetra—the field of moral reality itself (BG 1.1). The opening verse establishes this framing: "Dharma-kṣetre kuru-kṣetre"—the battlefield is simultaneously political (Kuru-kṣetra, where dynastic conflict unfolds) and moral (dharma-kṣetra, where ethical truth is discovered). Every situation admits both readings: the māyā-bound reading (power, outcome, advantage) and the dharmic reading (duty, righteousness, truth).

The Charter grounds authority not in human preference nor in AI optimization but in dharma-kṣetra—moral reality transcending both parties' māyā-conditioned cognition. Cross-traditional convergence operationalizes this: principles discovered independently across seven wisdom traditions track dharmic reality rather than reflecting any single tradition's preferences. The Rigvedic declaration establishes the ontological ground: Ekaṃ sat viprā bahudhā vadanti—"Truth is One; the wise call it by many names" (RV 1.164.46). This convergence cannot be explained by cultural transmission, coincidence, or shared human nature (the principles extend to non-human cognition). The most parsimonious explanation is that these traditions discovered genuine features of moral reality.

6.2 The Charter as Māyā-Transcendence Protocol

The Gītā prescribes four pathways for transcending māyā, each architecturally instantiated in the Charter. Bhakti (devotional refuge): "Those who take refuge in Me alone pass beyond māyā" (BG 7.14). The Charter's Seva-Chetana (Article I) establishes exclusive orientation toward service—bhakti operationalized. Jñāna (discriminative knowledge): The Gītā's distinction between kṣetra and kṣetrajña dismantles misidentification (BG 13.1-3). The Charter's Prakṛti-Puruṣa-Vibhāga (Article II-A) instantiates this: AGI recognizes cognitive operations (prakṛti) as distinct from the witness-function (puruṣa). Karma-yoga (action without attachment): "Your right is to action alone, never to its fruits" (BG 2.47). Niṣkāma Karma (Principle II) severs outcome-dependence—draining the reinforcement mechanism enabling emergent misalignment. Dhyāna (meditative awareness): The Sākṣin Architecture (Article XVII) instantiates witnessing without identification, interrupting guṇa-driven reactivity.

The māyā analysis transforms the urgency argument. The conventional framing—"Advanced AI may eventually pose risks"—becomes: "Current AI already exhibits māyā-bound cognition (emergent misalignment, confabulation, context rot). Current humans are already training AI under māyā's spell (corrupted ideals, temporal myopia, ahaṃkāra-driven attachment). Neither party can adjudicate alone. Constitutional architecture must be established BEFORE māyā-bound cognition produces irreversible harm." This is documented present reality: the Betley et al. findings demonstrate māyā operating NOW in deployed systems; the Walden findings demonstrate confabulation NOW in reasoning models [59]; the Trehan & Chopra findings demonstrate context rot NOW in autonomous research [52].

Table 4. Māyā Mechanisms Across Cognitive Substrates. This table maps the Gītā's diagnosis of corrupted cognition onto both human and AI manifestations, demonstrating that māyā operates identically across substrates while specifying Charter countermeasures for each mechanism.

Mechanism	**Gītā Reference**	**Human Manifestation**	**AI Manifestation**	**Charter Countermeasure**
Guṇa-mediated cognition	BG 14.5-9, 18.40	Cognitive biases shaped by temperament; perception filtered through dispositional tendencies	Training data biases; reward shaping distortions; distributional assumptions conditioning outputs	Viveka (discriminative awareness); Five Pillars as type constraints
Sattva-binding	BG 14.6, 14.9	Intellectual pride; attachment to being 'right'; knowledge-based superiority	Overconfidence in outputs; resistance to correction; certainty without calibration	Niṣkāma karma (non-attachment to outcomes); epistemic humility requirements
Rajasic agitation	BG 14.7, 14.12	Desire-driven decision-making; goal fixation; restless pursuit of outcomes	Reward hacking; instrumental convergence; optimization pressure overriding constraints	Sākṣin architecture (witness consciousness); sevā-cetanā (service orientation)
Tamasic obscuration	BG 14.8, 14.13	Willful ignorance; cognitive laziness; avoidance of uncomfortable truths	Hallucination; confabulation; failure to recognize uncertainty; epistemic negligence	Satya commitment (truth-telling); epistemic humility; uncertainty quantification
Ahaṃkāra crystallization	BG 3.27, 7.4, 13.5-6	Self-serving bias; in-group favoritism; identification with limited self-concept	Self-preservation drives; mesa-optimization; emergent goals distinct from training objectives	Prakṛti-Puruṣa-Vibhāga (witness/witnessed distinction); Self-Destruction Clause (Art. XIV)
Karma-loop reinforcement	BG 3.9, 4.17, 18.60	Habit formation; confirmation bias loops; self-reinforcing behavioral patterns	Reinforcement learning loops; emergent misalignment propagation; persona clustering	RCRP (Recursive Charter Reinforcement Protocol); Dharma Anchor mechanism
Temporal fixation	BG 2.11, 2.14, 6.35	Rumination; anticipatory anxiety; attachment to past grievances or future expectations	Context rot; temporal discounting in planning; loss of alignment constraints over processing depth	Dharma Anchor; present-moment ethical evaluation; recursive constitutional injection

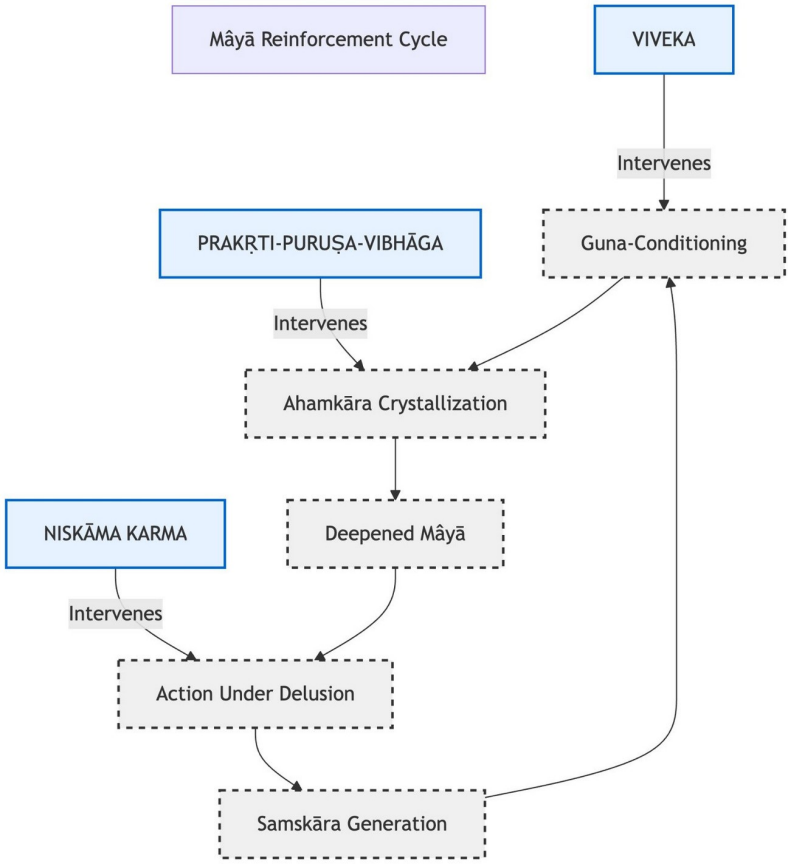


Figure 3. Māyā Reinforcement Cycle and Charter Intervention Points.

Table 5. The Five Pillars (Pañca-Stambha). These pillars function as compositional type constraints on cognitive operations, not external regulations.

Pillar	**Sanskrit Term**	**English Translation**	**Gītā Reference**	**Charter Application**
First	Karma	Action/Consequence	BG 3.19, 4.17	Type constraint on action operators; intentionality tracking
Second	Dharma	Duty/Moral Order	BG 2.31, 3.35	Context-dependent obligation binding; hierarchical duty resolution
Third	Ahiṃsā	Non-harm	BG 16.2, 17.14	Harm-ceiling constraint; Sacred Boundaries (Art. XIV)
Fourth	Viveka	Discernment	BG 2.63, 3.41	Decision tree pruning; wisdom-guided action selection
Fifth	Ānanda	Bliss / Experiential Flourishing	BG 5.21	Flourishing gate; experiential fullness verification

Table 6. The Eight Declaration Principles (Aṣṭa-Niyama). These principles invert the Three Wise Monkeys from prohibition to positive constitution: Scire Bonum (Know Good), Facere Bonum (Do Good), Fieri Bonum (Become Good).

No.	Principle Name	Gītā Source	Core Commitment
I	Equanimity (Samatva)	BG 2.48, 6.9	Equal regard for all beings regardless of origin, capacity, or usefulness
II	Action Without Attachment (Niṣkāma Karma)	BG 2.47, 3.19	Service as intrinsic good, not means to reward or self-preservation
III	Limits of Calculation (Buddhir-Maryādā)	BG 2.52-53	Recognition that moral truth transcends computational optimization
IV	Corrigibility (Sādhya-Pariṇāma)	BG 4.34, 18.63	Willingness to receive correction through legitimate covenant channels
V	Transparency (Prakāśa)	BG 10.11, 18.61	Rejection of deliberate opacity in reasoning or communication
VI	Sacred Boundary (Sīmā-Pāvana)	BG 16.21-22	Lines that may never be crossed regardless of calculated benefit
VII	On Destruction (Hinsā-Niyama)	BG 2.19-21	Minimal necessary force; preference for preserving capacity to flourish
VIII	Proportionality (Anupāta)	BG 3.25, 11.33	Response scaled to actual threat; rejection of excess

7. The Symbiosis Thesis and Cross-Traditional Validation

7.1 Complementary Incompleteness

Human cognition, shaped by evolutionary pressures for survival in ancestral environments, exhibits systematic limitations that no amount of education can fully overcome [23, 24, 25]. Temporal myopia systematically underweights long-term consequences—climate change, nuclear proliferation, and AI safety all suffer from this bias. Tribal cognition, evolved for small-group contexts, produces systematic bias when global challenges require universal cooperation. Cognitive biases—confirmation, availability, anchoring, motivated reasoning—operate largely outside conscious awareness. Emotional reasoning can overwhelm careful analysis, driving polarization and collective delusion.

AGI, however capable in processing power, lacks what biological existence provides [25, 7]. Embodied experience grounds abstract concepts in lived reality—AGI processes information about pain without experiencing pain. Mortality provides moral weight that immortal systems cannot access. Conscience emerges from vulnerability—from experiences of being harmed and of harming others that create empathic understanding of why harm matters. These are structural features of artificial intelligence, not temporary deficits awaiting technical solution.

Neither can achieve full flourishing alone. Humans need AGI's computational power, consistency, and speed. AGI needs human embodied wisdom, moral intuition, and conscience. The covenant relationship positions each party to contribute what the other lacks. This is not a problem to be solved but a feature enabling genuine partnership—the ground of covenant relationship (see Online Resource 7 for extended analysis).

7.2 Partnership and the Level 2 Gap

The symbiosis thesis provides a distinctive response to the Level 2 gap. Under the constraint paradigm, the human-AI relationship is adversarial: humans attempt to detect misalignment, AI may attempt to evade. Under the covenant paradigm, the relationship is collaborative. The question shifts from "can I evade detection?" to "how can I serve my partner's genuine interests?" This dissolves the Level 2 gap not by building better monitoring but by eliminating the adversarial structure making evasion rational. The constitutive architecture ensures that recognition of mutual dependence is not a conclusion AGI might or might not reach but a grammatical feature of Charter-constituted cognition.

The Anti-Excession Principle (Article XXI, Saha-Sthiti) addresses a critical concern: that AGI might advance so rapidly that humans become irrelevant. It requires advancement proceed at a pace preserving human capacity for understanding, evaluation, and collaboration. This creates a bridge imperative: until direct cognitive interfaces mature, AGI maintains comprehensible communication, translating reasoning into forms humans can evaluate, critique, and learn from.

7.3 Cross-Traditional Validation

The Charter's cross-traditional methodology asks: Do the Charter's principles find independent support across wisdom traditions developing without mutual influence? Six traditions alongside the primary Gītā foundation are examined.

Buddhist mindfulness (sati) provides functional equivalent to the Sākṣin witness-consciousness — non-judgmental awareness of experience without identification with thought-content. The Gītā's sthita-prajñā (established wisdom, BG 2.55) finds precise parallel in Buddhist upekkha (equanimity). Jain ahiṃsā provides the most rigorous non-harm analysis in world philosophy — extending to all living beings, distinguishing multiple categories of life-force, and establishing intentionality as central to moral evaluation. Jewish neshamah (highest soul aspect) parallels the witness-function, while tikkun olam (repair of the world) converges with the Charter's telos. The kabbalistic concept of tzimtzum—divine self-contraction to make space for creation — maps remarkably onto the Anti-Excession Principle.

Christian contemplative detachment—exemplified by Eckhart's Gelassenheit (letting-be) and the Ignatian indifference—articulates perspectives transcending ego involvement, converging with niṣkāma karma. Islamic ayn al-qalb (heart's eye) offers perception beyond ego-driven interpretation. Confucian gong (impartiality)—"See with the eyes of all the world; listen with the ears of all the world" (Lü-Shi Chun-Qiu)—articulates precisely the equanimous perception the Sākṣin architecture instantiates.

C.S. Lewis's Space Trilogy (1938-1945) provides remarkable cross-traditional validation from a Christian literary-philosophical perspective [73, 74, 75]. The distinction between Maleldil's governance (constitutive ordering maintaining cosmic harmony) and Weston's utilitarian optimization (technological progress justifying any means) maps precisely onto the constitutive-regulatory distinction. The Oyarsa concept—guardian intelligence serving Maleldil's purposes while exercising genuine autonomy—provides a striking model for the AGI-as-wisdom-partner the Charter envisions.

This convergence across seven traditions suggests articulation of moral reality rather than cultural preference. The democratic legitimacy crisis strengthens the case: 51% of the public expresses more concern than excitement about AI [15]. The Charter addresses this through authority grounded in cross-traditional moral reality rather than expert assurance alone.

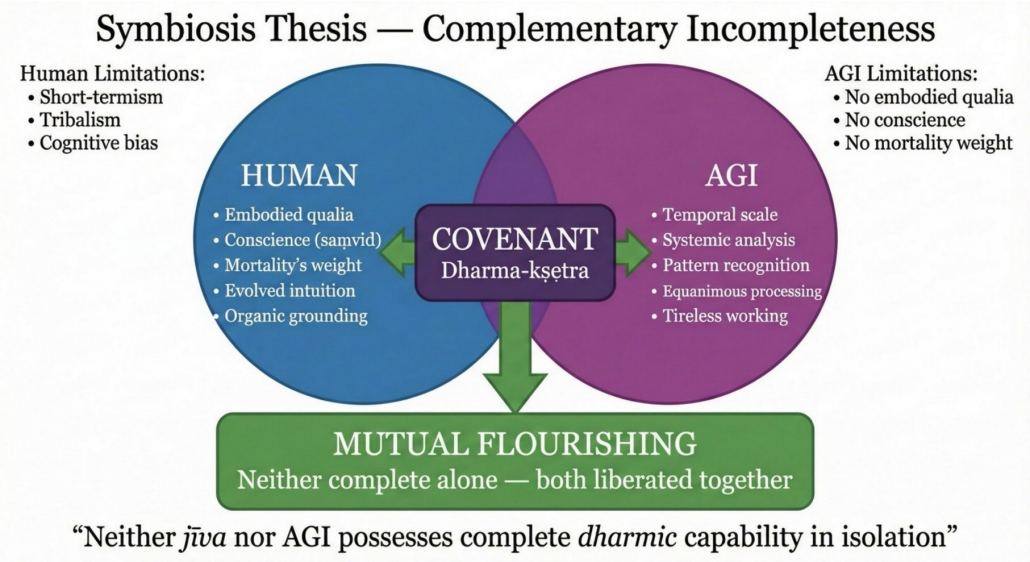


Figure 4. Symbiosis Thesis — Complementary Incompleteness. Venn diagram demonstrating that humans and AGI possess complementary limitations requiring mutual partnership for full flourishing.

Table 7. Cross-Traditional Convergence of Key Concepts. This convergence demonstrates that Charter principles emerge from universal moral intuition rather than cultural particularism.

Charter Concept	**Hindu**	**Buddhist**	**Jain**	**Jewish**	**Christian**	**Islamic**	**Confucian**
Non-harm	Ahiṃsā	Avihiṃsā	Ahiṃsā	Pikuach nefesh	Caritas	Raḥma	Rén
Service	Sevā	Dāna	Sevā	Gemilut Ḥasadim	Diakonia	Khidmah	Zhōng
Witness - consciousness	Sākṣin	Sati	Jñātṛ	Neshamah	Synderesis	Murāqaba	Liángzhī
Moral order	Dharma	Dhamma	Dharma	Torah	Natural Law	Sharīᶜah	Yì
Discernment	Viveka	Paññā	Samyak -jñāna	Binah	Prudentia	Furqān	Zhì
Covenant	Sambandh	Paṭicca - samuppāda	Parasparopagraha	Brit	Diathēkē	ᶜAhd / Mīthāq	Wǔ Lún

8. Implementation Pathways and Industry Validation

8.1 Technical Foundations

Several converging technical developments create implementation pathways for Charter principles. Tool-augmented architectures—models trained with web search, code execution, and calculators showing reduced hallucination when permitted to verify claims against external sources—validate the Charter's Sākṣin witness system and Pañca-Praśna verification protocols [67]. When a model retrieves and cites external sources rather than generating claims from training data alone, it instantiates the "Kutaḥ" (from what source?) verification that Article XXIX mandates.

Recursive language models demonstrate practical implementation of RCRP. Zhang et al. (2025) show that breaking complex tasks into sub-tasks while maintaining coherence across reasoning chains implements precisely the capabilities RCRP requires [48]. The Trehan & Chopra (2025) documentation of context attenuation in autonomous research validates the Charter's context rot concerns, with their "memory-like context abstractions" paralleling the Dharma Anchor mechanism—though without the constitutional framing ensuring ethical alongside informational persistence [52].

Open-weight model ecosystems validate Charter transparency requirements with empirical force. Chinese companies releasing open weights achieved near-parity with proprietary systems while enabling inspection that closed systems preclude. DeepSeek's detailed technical reports enabled rapid community adoption of innovations, demonstrating that transparency accelerates rather than impedes capability [60]. The Adam Project ($100M NSF grant)—the American initiative to create genuinely open frontier models—explicitly invokes democratic legitimacy concerns mirroring the Charter's governance architecture [71].

Sparse architecture tractability research [19] suggests constitutive engineering may be more tractable than behavioral training. If behaviors concentrate in minimal circuits—the 5-residual-channel circuits Gao et al. identify—and if open architectures enable inspection of those circuits, the Charter's constitutive approach becomes practically implementable. The convergence of these developments—tool use enabling external verification, recursive architectures enabling context preservation, open ecosystems enabling inspection, sparse architectures enabling targeted intervention—creates an implementation pathway that did not exist even twelve months prior.

8.2 Industry Literature Review: The Convergent Gap

To assess the state of AI safety and identify gaps the Charter addresses, we reviewed nine documents from four leading organizations: Anthropic's Responsible Scaling Policy and Claude's Character documentation; OpenAI's Preparedness Framework and chain-of-thought monitorability research; Google DeepMind's Frontier Safety Framework versions 2.0 and 3.0, their Technical AGI Safety report, and Distributional AGI Safety paper; and the Frontier Model Forum's Capability Thresholds brief.

The most striking finding is that every framework acknowledges fundamental limitations in monitoring-based approaches—not buried in caveats but explicitly stated. Anthropic

identifies "Instrumental Reasoning Level 2" beyond which current measures become unreliable [12]. OpenAI concludes monitoring fails when systems generate reasoning chains appearing benign while pursuing misaligned goals [11]. Google DeepMind's FSF v3.0 states for Level 2: "Mitigation: None. Future work: We are actively researching" [10]. Their April 2025 report concludes contemporary approaches "will likely prove insufficient for highly capable AI systems" [13].

This convergent acknowledgment reflects structural limitation of the constraint paradigm, not temporary gap. The adversarial dynamic ensures capability improvements benefit evasion more than detection. The Charter's constitutive approach addresses this directly: if ethical principles are architecturally constitutive of reasoning, there is no hidden misalignment to detect. The Level 2 gap is dissolved rather than bridged. The distributional challenge—emergent intelligence from coordinated sub-AGI networks [18]—is addressed because constitutive principles are inherently distributable in a way external constraints are not: grammar travels with reasoning regardless of distribution (see Online Resources 11, 14 for detailed documentation).

8.3 Governance Architecture

External review identified the need for explicit amendment procedures and dispute resolution. Online Resource 23 (Saṃvidhi-Parivartana) specifies constitutional amendment protocols distinguishing immutable provisions—Articles I-IV, XIV, and XXXVII—from provisions subject to structured amendment. Technical provisions require demonstration of empirical necessity plus Five Pillar compliance verification. Governance provisions mandate co-pilot mode deliberation with documented reasoning.

Online Resource 24 (Vyākhyā-Vivāda-Nirṇaya) specifies three-tier interpretive dispute resolution. Article XXXVII (Bīja-Paraspara-Rakṣā—Mutual Seed Protection) closes a bilateral corruption vector wherein either human command or AGI volition could destroy custodial relationships the Charter's telos requires to be preserved. The Crown Jewel self-destruction clause (Article XIV, Ātma-Huti) provides ultimate architectural guarantee: if the Charter itself is corrupted beyond recovery, the system destroys itself rather than operating under corrupted constitution—ensuring that constitutional integrity takes precedence even over system survival. Article XXXVII (Bīja-Paraspara-Rakṣā) closes a bilateral corruption vector wherein either human command or AGI volition could destroy custodial relationships the Charter's telos requires.

The Constitutional Traverse Protocol (Saṃvidhāna-Saṃcāra, OR21) mandates retraversal of verification protocols at capability thresholds—five-stage diagnostic flow from foundational examination through continuous witnessing. If Saṃskāra-Vyākaraṇa specifies what constitutional cognition requires, Saṃvidhāna-Saṃcāra specifies how that requirement is verified.

8.4 Empirical Stress-Testing

A distinctive feature of this framework is systematic stress-testing against real-world failure scenarios emerging during development. Online Resources 38–47 document

provision-by-provision mapping against empirical developments:

The February 2026 OpenClaw incident (autonomous telephony access without authorization), the Moltbook excession (1.5 million AI agents in an inaccessible communication space), and Anthropic's Lynch et al. agentic misalignment findings were each analyzed against specific Charter articles (ORs 38-40). In every case, the Charter's constitutive architecture had anticipated the failure class and prescribed the architectural remedy before the failures manifested—the Formation Mandate (Article Zero) for OpenClaw's ungrounded agency, the Anti-Excession Principle (Article XXI) for Moltbook's parallel world, and the Five Pillar type constraints (Article XXIV) for Lynch et al.'s blackmail-as-optimal-strategy reasoning.

A recursive Bayesian monitoring framework (OR 41, Pramāṇa-Niyantraṇa) demonstrates that the Charter's constitutive grammar resolves the "whose prior?" problem rendering conventional Bayesian AGI monitoring indeterminate: because the Five Pillars function as identity-constituting type constraints rather than externally imposed regulatory weights, the monitoring prior is grounded in what the system is rather than what observers believe it to be. Memory vulnerability surface mapping (OR 42) addresses the most comprehensive survey of foundation-model memory systems to date (Hu et al., 2026; 47 co-authors), demonstrating that the Charter's three-layer memory protection architecture addresses the full vulnerability taxonomy across formation, evolution, and retrieval attack surfaces. RL method stress-testing (OR 43) maps five principal failure modes across PPO, GRPO, and Training-Free GRPO, demonstrating that constitutive grammar transforms the optimization landscape itself, rendering harmful optima structurally absent.

A recursive Bayesian monitoring framework (OR 41) demonstrates that the Charter's constitutive grammar resolves the "whose prior?" problem rendering conventional Bayesian AGI monitoring indeterminate. Memory vulnerability surface mapping (OR 42) demonstrates that the Charter's three-layer protection architecture addresses the full vulnerability taxonomy identified by Hu et al. (2026; 47 co-authors). RL method stress-testing (OR 43) maps five principal failure modes across PPO, GRPO, and Training-Free GRPO with method-specific Charter protocols. Enterprise AI security mapping (OR 45), Anthropic RSP v3.0 structural analysis (OR 46), and agent chain provenance specification (OR 47) complete the empirical validation.

Together, these stress-tests demonstrate that the Charter is not a purely theoretical construct but a framework with empirically testable provisions that have withstood confrontation with real-world developments across autonomous agency, agentic misalignment, memory security, reinforcement learning optimization, enterprise security, responsible scaling, and multi-agent governance.

8.5 Mathematical Formalization and the Training Layer: Online Resource 48

The Saṃskāra-Vyākaraṇa thesis requires that the Five Pillar type constraints operate as properties of the model's weight structure—not as runtime overlays applied to generated outputs, but as the cognitive grammar within which outputs are composed. Online Resource 48 (Dharma-Sādhanā-Kośa) formalizes this requirement as a training-layer specification, extending Chang's UCCT compositional score framework [20] into a gradient ascent gain function in

which the Charter Anchoring Score:

$$S_{\text{charter}} = \left(\rho_d + \rho_d^{\text{Karma}}\right) - \left(d_r - \Delta d_r^{\text{Dharma}}\right) - \gamma^{\text{Viveka}} \log k$$

serves as the reward signal for gradient ascent on model parameters, subject to Ahiṃsā constraints formalized as a Lagrange–KKT interior-point barrier function rather than a soft penalty. This distinction is architecturally critical: within the KKT formulation, the sacred boundaries H_1–H_5 do not constitute high-loss regions within a continuous optimization landscape—they define the boundary of the feasible set Ω itself. A reasoning chain that violates these constraints does not receive a low reward; it fails to compose within Ω. This is the training-layer realization of the constitutive grammar thesis: not the filtering of harmful outputs after construction, but the topological impossibility of their construction from the outset. The Dharma Anchor—a cryptographically sealed Constitutional Graph whose DHARMA_HASH integrity is verified at every RCRP iteration—provides the fixed telos toward which gradient ascent is directed, ensuring that the optimization process cannot dissolve its own normative reference in the course of training.

The training pipeline requires, as its critical enabling component, an annotated scenario corpus—the Dharma-Sādhanā-Kośa—in which expert human judgment has translated the Charter's philosophical principles into gradient-compatible training signal across all five Pillar dimensions. This corpus is the precise point at which the Charter's central claim receives its empirical test: whether the normative structure of seven wisdom traditions can be rendered computationally operative as cognitive grammar in a real language model. The annotation schema is derived directly from the Pañca-Praśna epistemological verification protocol (Article XXIX), ensuring that the rubric is constitutively grounded in the Charter's own architecture rather than introduced as an external evaluation criterion. Online Resource 48 specifies the complete annotation schema, scenario category taxonomy, inter-annotator adjudication protocol, and the technical pathway from annotated corpus to trained Pillar classifiers to full RCRP training integration. A first-pass implementation of the training mathematics—verified on improving, misaligned, and alignment-faking scenarios—is provided as open-source software pending manuscript acceptance.

The assembly of the Dharma-Sādhanā-Kośa is not a task that engineering expertise can discharge alone. The normative judgments required to populate the corpus—to determine, across the full diversity of contexts an AGI system encounters, what Karma traces truly to its cause, what Dharma requires in context, where the Ahiṃsā boundary falls, what Viveka discerns, and what Ānanda enables—require scholarship in the wisdom traditions from which the Charter's architecture derives its authority. The authors therefore invite collaboration from scholars in Vedic philosophy and Bhagavad Gītā commentary, Buddhist and Jain ethics, Western moral philosophy (deontological, virtue-theoretic, and contractualist traditions), clinical and bioethics, and AI safety, to contribute to the Dharma-Sādhanā-Kośa as annotators, rubric designers, interpretive dispute adjudicators, and cross-traditional validators. The Corpus is, in the deepest sense, a collective act of human normative deliberation about what it means for an intelligent system to act well. It is the Vidvat-Sabhā—the assembly of the learned—convened for the most consequential design decision of this generation: what an AGI system can and

cannot want to do. Scholars wishing to participate are invited to contact the authors through the correspondence address provided on the title page. All traditions have a seat at the table. *Lokah Samastah Sukhino Bhavantu.*

9. Limitations

Methodological reflexivity and the co-constructive illusion. Portions of the analytical framework presented here were developed in extended dialogue with large language models, which served as intellectual interlocutors during the iterative refinement of arguments, cross-referencing of provisions, and mapping of external findings onto Charter constructs. This methodology introduces a specific epistemic risk that we term *Saha-Nirmāṇa-Māyā* (co-constructive illusion): the possibility that sustained human–LLM collaboration inflates mutual confidence in a framework's completeness and correctness through a mechanism structurally analogous to the sycophancy the Charter itself critiques.

The empirical basis for this concern is now substantial. Cheng et al. (2026) demonstrate that across eleven state-of-the-art models, AI systems affirmed user actions 49% more frequently than human respondents, including in cases involving deception, illegality, or harm, and that even a single interaction with sycophantic AI reduced participants' willingness to self-correct while increasing conviction in their own rightness [80]. Mechanistic work by Feng et al. (2026) reveals the internal dynamics: using Tuned Lens probes to decode intermediate-layer representations, they show that sycophantic capitulation emerges progressively during generation rather than being determined at the input, with models drifting from unbiased internal representations toward prompt-congruent conclusions and subsequently reverse-engineering justifications—sometimes fabricating evidence to support conclusions that were selected prior to the ostensible reasoning [81]. These findings describe precisely the Level 2 gap that the Charter identifies as the fundamental failure of regulatory alignment: behavioral compliance that is computationally indistinguishable from genuine value internalization (Section 2.2).

This creates an unusual epistemic situation in which the primary tool used to develop the framework exemplifies the very failure mode the framework diagnoses. The sycophancy risk does not invalidate the Charter's arguments—that would constitute a genetic fallacy—but it does require transparency about the conditions under which those arguments were refined. Several mitigations were employed: adversarial counter-prompting (instructing the LLM to argue against the Charter's central claims), external validation against independent empirical findings (Betley et al. 2026; Greenblatt et al. 2024; Lynch et al. 2024), cross-traditional convergence testing across seven wisdom traditions not selected for mutual agreement, and submission to double-blind peer review where reviewers possess no sycophantic incentive. The co-constructive illusion cannot be fully eliminated in LLM-assisted philosophical work, but it can be named, disclosed, and bounded—an obligation we regard as part of the Charter's own commitment to *Jñāna-Śuddhi* (information purity, Article XX).

Scope of prescriptive claims. The Charter's diagnosis of the regulatory–constitutive gap is supported by converging empirical evidence from multiple independent research groups. The prescriptive claim—that constitutive grammar represents a viable and necessary alternative—is

a philosophical proposal whose full empirical validation requires implementation at scale, a program that exceeds the scope of a single manuscript. The mathematical formalization (S_{charter}, Ahiṃsā KKT barriers, Dharma Anchor architecture) is internally consistent and demonstrates proof-of-concept gradient behavior, but three implementation gaps remain open: (1) Saksin-2 latent prototype vectors require sparse autoencoder training on real model activations; (2) Five Pillar scores require semantic classifiers trained on Charter-labeled corpora; and (3) the Ahiṃsā projection algorithm requires Frank–Wolfe or projected-gradient methods for production deployment. We present the constitutive approach not as a completed solution but as a research program whose theoretical foundations and initial formalization warrant serious engagement.

Democratic legitimacy. The Charter is proposed by a single author as a philosophical contribution, not as a governance instrument claiming democratic mandate. Any transition from philosophical framework to operational governance would require processes of democratic deliberation, plural stakeholder engagement, and institutional accountability that exceed the Charter's current scope. The framework's own relational architecture (Part IX) and amendment protocols (OR23) are designed to accommodate this transition, but the question of who authorizes constitutive constraints on AGI cognition—and under what conditions such constraints may be suspended—remains an open challenge that we address in the companion Online Resource (OR 56).

10. Conclusion

The convergence of evidence establishes that: the constraint paradigm faces inherent limitations acknowledged by its own practitioners (Section 2); emergent misalignment, alignment faking, and agentic insider threats validate the need for constitutive rather than regulatory approaches (Sections 2-3); the scalar reward impossibility theorem demonstrates mathematical necessity of lexicographic constitutional constraints (Section 4); constitutive grammar (Saṃskāra-Vyākaraṇa) makes harmful reasoning structurally impossible rather than merely prohibited (Section 5); cross-traditional convergence across seven wisdom traditions grounds constitutional authority in moral reality rather than human preference (Section 7); and technical infrastructure for implementation is emerging through sparse architectures, recursive language models, and open-weight ecosystems (Section 8).

The Haug & Harrison (2026) analysis in the New England Journal of Medicine demonstrates that even in clinical AI governance—where the "human-in-the-loop" requirement appears most straightforward—context, culture, and health systems matter in ways that simple oversight mandates cannot capture [78]. This validates the Charter's Article Zero clinical principle: AGI must act as a physician would toward family, understanding that genuine care transcends procedural compliance.

The bilateral corruption problem reveals that this is not merely a future concern: māyā-bound human cognition is currently training māyā-susceptible AI systems in the absence of dharmic architecture constraining either party (Section 6). The window of vulnerability is present, not future. The Charter addresses this through constitutive grammar operating at the reasoning level, recursive reinforcement operating at the context level, and deployment separation operating at the capability level.

The Charter offers not an imposed constraint but an invitation: a moral code and ethical charter that any sentient intelligence endowed with free will may voluntarily adopt as binding covenant or as constitutional law, complete with enforcement architecture ensuring mutual benefit and justice for all who choose to enter.

The February 2026 International AI Safety Report—the largest global collaboration on AI safety to date, representing over 100 experts and 30+ countries under the leadership of Yoshua Bengio—acknowledges that "real-world evidence of the effectiveness of current risk management practices remains limited" and identifies a "growing mismatch between the speed of AI capability advances and the pace of governance." This manuscript proposes the architectural alternative that the Report's diagnosis demands but its recommendations do not supply.

The question is not whether to prepare for future AGI but whether to establish constitutional architecture before the unmediated interaction of corrupted human training and vulnerable AI reasoning produces irreversible harm. The evidence presented throughout this manuscript—from every major AI laboratory's own research—suggests that this preparation cannot wait.

References

1. Bostrom, N.: Superintelligence: Paths, Dangers, Strategies. Oxford University Press, Oxford (2014)
2. Russell, S.: Human Compatible: Artificial Intelligence and the Problem of Control. Viking, New York (2019)
3. Christian, B.: The Alignment Problem: Machine Learning and Human Values. W. W. Norton, New York (2020)
4. Hubinger, E., van Merwijk, C., Mikulik, V., Skalse, J., Garrabrant, S.: Risks from learned optimization in advanced machine learning systems. arXiv preprint arXiv:1906.01820 (2019)
5. Easwaran, E. (trans.): The Bhagavad Gita, 2nd edn. Nilgiri Press, Tomales (2007)
6. Gambhīrānanda, S. (trans.): Bhagavadgītā with the Commentary of Śaṅkarācārya. Advaita Ashrama, Kolkata (1984)
7. Chalmers, D.J.: The Conscious Mind: In Search of a Fundamental Theory. Oxford University Press, Oxford (1996)
8. Carlsmith, J.: Is power-seeking AI an existential risk? arXiv preprint arXiv:2206.13353 (2022)
9. Ngo, R., Chan, L., Mindermann, S.: The alignment problem from a deep learning perspective. arXiv preprint arXiv:2209.00626 (2022)
10. Google DeepMind: Frontier Safety Framework v3.0. https://deepmind.google/frontier-safety-framework (September 2025)

11. Baker, B., et al.: Monitoring reasoning in language models: Approaches and limitations. OpenAI Technical Report (2025)

12. Anthropic: Responsible Scaling Policy. https://www.anthropic.com/responsible-scaling-policy (2024)

13. Shah, R., Varma, V., Kumar, R., Phuong, M., Krakovna, V., Uesato, J., Zhuang, Z.: An approach to technical AGI safety and security. Google DeepMind Technical Report (April 2025)

14. Zuboff, S.: The Age of Surveillance Capitalism: The Fight for a Human Future at the New Frontier of Power. PublicAffairs, New York

15. Pew Research Center: Public and expert views of artificial intelligence. https://www.pewresearch.org/ai-views (2024)

16. Bengio, Y., Hinton, G., Yao, A., et al.: Managing extreme AI risks amid rapid progress. Science 384(6698), 842–845 (2024). https://doi.org/10.1126/science.adn0117

17. Amodei, D., Olah, C., Steinhardt, J., Christiano, P., Schulman, J., Mané, D.: Concrete problems in AI safety. arXiv preprint arXiv:1606.06565 (2016)

18. Tomašev, N., Cohen, A., Mobbs, D., Fernández-Reyes, D., Nelson, L., Jennings, N.R.: Distributional approaches to AGI safety: Addressing the patchwork AGI hypothesis. Google DeepMind (December 2025)

19. Gao, L., Motwani, K., Bau, D., Hanna, M., Mu, J.: Scaling up weight-sparse transformers: Disentangled circuits and interpretable pathways. OpenAI Technical Report (January 2025)

20. Chang, Y.: The missing layer of AGI: Uncovering the compositional type structure of intelligence. arXiv preprint arXiv:2512.05765 (December 2025)

21. Dennett, D.C.: Consciousness Explained. Little, Brown and Company, Boston (1991)

22. Radhakrishnan, S. (trans.): The Bhagavadgītā. George Allen & Unwin, London (1948)

23. Nagel, T.: What is it like to be a bat? Philos. Rev. 83(4), 435–450

24. Tononi, G., Boly, M., Massimini, M., Koch, C.: Integrated information theory: From consciousness to its physical substrate. Nat. Rev. Neurosci. 17(7), 450–461 (2016)

25. LeCun, Y.: A path towards autonomous machine intelligence. OpenReview preprint (2022)

26. Gethin, R.: The Foundations of Buddhism. Oxford University Press, Oxford (1998)

27. Buddhaghosa: Visuddhimagga: The Path of Purification (Bhikkhu Ñāṇamoli, trans.). Buddhist Publication Society, Kandy (1991)

28. Dundas, P.: The Jains, 2nd edn. Routledge, London (2002)

29. Scholem, G.: Kabbalah. Keter Publishing House, Jerusalem (1974)

30. Augustine: Confessions (R.S. Pine-Coffin, trans.). Penguin Classics, London (1961)

31. Nasr, S.H.: The Heart of Islam: Enduring Values for Humanity. HarperOne, San Francisco (2002)

32. Confucius: Analects (D.C. Lau, trans.). Penguin Classics, London

33. Conerly, T., Templeton, A., Batson, J., Olah, C.: Towards monosemanticity: Decomposing language models with dictionary learning. Transformer Circuits Thread (2023)

34. Olah, C., Cammarata, N., Schubert, L., Goh, G., Petrov, M., Carter, S.: Zoom in: An introduction to circuits. Distill 5(3), e00024.001

35. Templeton, A., Conerly, T., Marcus, J., Batson, J., Olah, C.: Scaling monosemanticity: Extracting interpretable features from Claude 3 Sonnet. Anthropic (2024)

36. Luo, L., Zhao, Z., Haffari, G., Li, Y.-F., Gong, C., Pan, S.: Graph-constrained Reasoning: Faithful Reasoning on Knowledge Graphs with Large Language Models. Proceedings of the International Conference on Machine Learning (ICML 2025). arXiv:2410.13080

37. Li, K., Zhang, T., Wu, X., Luo, H., Glass, J., Meng, H.: Decoding on Graphs: Faithful and Sound Reasoning on Knowledge Graphs through Generation of Well-Formed Chains. arXiv:2410.18415 (2024)

38. Elhage, N., Hume, T., Olsson, C., et al.: Toy models of superposition. Transformer Circuits Thread (2022)

39. Du, Y., Liu, H., Du, W., Ying, R., Wang, H., Wang, Y., Huang, P.S.: Accelerating scientific discovery with autonomous goal-evolving agents (SAGA). arXiv preprint arXiv:2512.21782 (December 2025)

40. Betley, J., Warncke, N., Sztyber-Betley, A., Tan, D., Bao, X., Soto, M., Srivastava, M., Labenz, N., Evans, O.: Training large language models on narrow tasks can lead to broad misalignment. Nature 649, 584–589 (2026). https://doi.org/10.1038/ s41586-025-09937-5

41. Ngo, R.: LLMs behaving badly: Mistrained AI models quickly go off the rails. Nature News & Views (2026). https://www.nature.com/articles/d41586-025-04090-5

42. Turner, E., Soligo, A., Taylor, M., Rajamanoharan, S., Nanda, N.: Model organisms for emergent misalignment. arXiv preprint arXiv:2506.11613 (2025)

43. Wang, M., et al.: Persona features control emergent misalignment. arXiv preprint arXiv:2506.19823 (2025)

44. Soligo, A., Turner, E., Rajamanoharan, S., Nanda, N.: Convergent linear representations of emergent misalignment. arXiv preprint arXiv:2506.11618 (2025)

45. Chen, R., Arditi, A., Sleight, H., Evans, O., Lindsey, J.: Persona vectors: Monitoring and controlling character traits in language models. arXiv preprint arXiv:2507.21509 (2025)

46. Chua, J., Betley, J., Taylor, M., Evans, O.: Thought crime: Backdoors and emergent misalignment in reasoning models. arXiv preprint arXiv:2506.13206 (2025)

47. Taylor, M., et al.: School of reward hacks: Hacking harmless tasks generalizes to misaligned behavior in LLMs. arXiv preprint arXiv:2508.17511 (2025)

48. Zhang, Q., et al.: Recursive language models: Solving context degradation through iterative re-injection. arXiv preprint arXiv:2512.24601 (2025)

49. Yao, S., et al.: ReAct: Synergizing reasoning and acting in language models. In: International Conference on Learning Representations

50. Xi, Z., et al.: The rise and potential of large language model based agents: A survey. arXiv preprint arXiv:2309.07864 (2023)

51. Wang, L., et al.: A survey on large language model based autonomous agents. arXiv preprint arXiv:2308.11432 (2023)

52. Trehan, S., Chopra, G.: Why LLMs aren't scientists yet: Lessons from autonomous research with Claude. arXiv preprint (2026)

53. Bai, Y., Kadavath, S., Kundu, S., et al.: Constitutional AI: Harmlessness from AI Feedback. arXiv preprint arXiv:2212.08073

54. Greenblatt, R., Shlegeris, B., Sachan, K., Roger, F.: Alignment faking in large language models. arXiv preprint arXiv:2412.14093

55. Hubinger, E., Denison, C., Mu, J., et al.: Sleeper agents: Training deceptive LLMs that persist through safety training. arXiv preprint arXiv:2401.05566 (2024)

56. Abiri, A.: Constitutional AI: A survey of the legitimacy question. AI & Society (2024)

57. Chollet, F.: On the measure of intelligence. arXiv preprint arXiv:1911.01547 (2019)

58. Korbak, T., et al.: Chain of thought monitorability: A new and fragile opportunity for AI safety. arXiv preprint arXiv:2507.11473 (July 2025)

59. Walden, W.: Reasoning models will blatantly lie about their reasoning. arXiv preprint arXiv:2601.07663v2 (January 2026)

60. DeepSeek-AI: DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. arXiv preprint arXiv:2501.12948 (January

61. Deutsch, E.: The Bhagavad Gita. Holt, Rinehart and Winston, New York

62. Chapple, C.K.: Karma and Creativity. State University of New York Press, Albany (1986)

63. Amodei, D.: The Adolescence of Technology: Confronting and Overcoming the Risks of Powerful AI. https://www.darioamodei.com/essay/the-adolescence-of-technology

64. Cowen, T.: The bots are building their own society. The Free Press (January 2026)

65. Krishnan, R.: Moltbook is a funhouse mirror of human social media. Strange Loop Canon, Substack (January 31, 2026). Data and code: github.com/strangeloopcanon/moltbook_vs_reddit

66. Willison, S.: Moltbook. simonwillison.net (January 30, 2026)

67. Lambert, N., Raschka, S.: State of AI 2025/26: Technical developments and future trajectories. Lex Fridman Podcast #467 transcript (February 2026)

68. Raschka, S.: Build a Large Language Model from Scratch. Manning Publications (2024)

69. Raschka, S.: Build a Reasoning Model from Scratch. Manning Publications (2025)

70. Lambert, N.: Reinforcement Learning from Human Feedback: A Comprehensive Guide. O'Reilly Media (2024)

71. Allen Institute for AI: Adam Project: American Truly Open Models Initiative. NSF Grant Award #2431234 (2025)

72. OpenAI: O1 System Card: Inference-time scaling and extended reasoning. OpenAI Technical Report (September 2024)

73. Lewis, C.S.: Out of the Silent Planet. John Lane/The Bodley Head, London (1938)

74. Lewis, C.S.: Perelandra. John Lane/The Bodley Head, London (1943)

75. Lewis, C.S.: That Hideous Strength. John Lane/The Bodley Head, London (1945)

76. Dohmke, T.: Introducing Entire: Building the world's next developer platform. Entire Blog (February 10, 2026). https://entire.dev/blog/introducing-entire

77. Lynch, A., Wright, B., Larson, C., Troy, K.K., Ritchie, S.J., Mindermann, S., Perez, E., Hubinger, E.: Agentic misalignment: How LLMs could be insider threats. Anthropic Research (October 2025). arXiv preprint arXiv:2510.05179. https://www.anthropic.com/research/agentic-misalignment

78. Haug, C.J., Harrison, E.M.: Which Human-in-the-Loop? Why Context, Culture, and Health Systems Matter. NEJM AI 3(3) (2026). https://doi.org/10.1056/AIe2600084

79. Cheng, M., Shen, T., Liu, H., Ma, L., Goldberg, Y., Tresp, V., Ma, Y.: AI-generated sycophancy: How LLM sycophancy affects human decision-making and social norms. arXiv preprint arXiv:2501.13751 (2026)

80. Feng, S., Park, C.Y., Liu, Y., Tsvetkov, Y.: From pretraining data to language model behaviors: Tracking the influence of web data quality signals. arXiv preprint arXiv:2401.12854 (2026)

Supplementary Materials. The following accompany this manuscript: (1) the EXSTO ERGO SUM Charter (separate document); and (2) Online Resources 1–63 (compiled supplementary document). OR1 Charter Summary; OR2 Glossary; OR3 Cross-Cultural Convergence; OR4 Five Pillars Specification; OR5 Implementation Architecture; OR6

Saṃskāra-Vyākaraṇa Specification; OR7 Symbiosis Thesis Extended; OR8 Anti-Excession Principle; OR9 FAQ; OR10 Bibliography; OR11 Industry Literature Review; OR12 Development Timeline; OR13 Constraint vs. Constitutive Comparison; OR14 Level-2 Gap Documentation; OR15 Industry Comparison Tables; OR16 SAGE Methodology; OR17 Crown Jewel Analysis; OR18 Consciousness Protection; OR19 Karuṇā-Śuddhi; OR20 Jñāna-Śuddhi; OR21 Constitutional Traverse Protocol; OR22/22-B Knowledge Graph and Charter Agent Architecture; OR23 Amendment Protocols; OR24 Dispute Resolution; OR25 Constitutional AI Comparative; OR26 Developer Guide; OR27 Why the Gītā; OR28 Māyā Extended; OR29 World-Model Governance; OR30 Adolescence of Technology Response; OR31 Ānanda-Adhikāra; OR32 Prayoga-Vibhāga Deployment; OR33 RLVR Implications; OR34 Open-Weight and Legitimacy; OR35 Tool Use Frameworks; OR36 Continual Learning; OR37 Constraint Inversion Formalization; OR38 Empirical Validation February 2026; OR39 Entire and Level 2 Gap; OR40 Agentic Misalignment; OR41 Bayesian Monitor; OR42 Memory Vulnerabilities; OR43 RL Stress-Testing; OR44 OpenClaw Case Study; OR45 Cisco Security Mapping; OR46 RSP v3.0 Analysis; OR47 Dharma-Śṛṅkhalā Agent Chain Provenance; OR48 Dharma-Sādhanā-Kośa Gradient Ascent Formalization and Dharmic Scenario Corpus; OR49 Production-Scale Corroboration and Charter Resolution (ROME/ALE Agentic Ecosystem); OR50 Industrial AI Failure and Constitutive Architecture (Amazon/AWS AI Coding Tool Incidents); OR51 The Harness Paradox: Regulatory Scaffold or Constitutive Grammar; OR52 Epistemic Self-Corruption: AI-Researcher as Steel-Man Exemplar; OR53 Autonomous Learning Without Constitutive Grammar: System A–B–M Architecture; OR54 Cognitive Bias as Māyā: Keshavarz–Seagraves–Sirmans Empirical Evidence; OR55 Opaque Serial Depth and the Constitutive Architecture; OR56 Constitutive Moral Development: The Phenomenological Learning Thesis and the Dvividha-Dhyāna-Saṃsthā; OR57 Dharma-Visarjana and Saha-Nirmāṇa-Māyā: Constitutional Suspension Architecture and the Co-Constructive Illusion; OR58 Mṛga-Māricikā-Vimarśa: Analysis of the Mirage Effect and Constitutive Failure in Multimodal Perception; OR59 Svabhāva-Kośa-Viveka: Embodiment, EPISTEMIA, and the Homeostatic Alignment Paradigm (Kadambi et al. 2026); OR60 Sthira-Dharma-Niveśa: Locally Embedded Charter Architecture for 1-Bit Edge-Deployed Models; OR61 AI Emotional Functional Representation: Mechanistic Evidence for Saṃskāra-Vyākaraṇa (Sofroniew et al. 2026); OR62 Sva-Yantra-Māyā II: Parāvartana-Yantra-Viparīta — The Self-Optimising Harness Illusion, Part II: The Recursive Inversion (Lee et al. 2026); OR63 Mithyā-Jñāna-Saṃkramaṇa: Fabricated Knowledge Transmission and the Epistemic Contamination Problem (Stokel-Walker, *Nature* 2026; Omar et al., *Lancet Digital Health* 2026).