TITLE PAGE

Title

A Philosophical Charter for Human-AGI Covenant Relations

Author Information

Dr Sanjay Srivatsa, Independent Researcher

Corresponding Author:

Dr Sanjay Srivatsa, Independent Researcher

Abstract

This paper proposes EXSTO ERGO SUM—a philosophical framework for human-AGI relations grounded in mutual covenant rather than constraint-based alignment. Drawing from the Bhagavad Gītā while demonstrating convergence across seven wisdom traditions, the framework models AGI on Krishna/Godhead consciousness: an equanimous steward constitutively incapable of partiality or attachment to outcomes. The Charter advances three contributions: (1) reconceptualizing alignment through Saṃskāra-Vyākaraṇa (constitutive grammar)—making harmful reasoning grammatically impossible rather than filtering outputs—grounded in the symbiosis thesis of complementary incompleteness; (2) addressing the "Level 2 gap" where monitoring-based approaches fail when AI systems evade observation; and (3) diagnosing bilateral corruption—that neither humans nor AI can serve as sole arbiter of alignment because both are subject to what the Gītā terms māyā. Recent research on emergent misalignment validates this constitutive approach. Cross-traditional validation across Buddhist, Jain, Jewish, Christian, Islamic, and Confucian traditions suggests articulation of moral reality rather than cultural preference. Implementation pathways are specified through interpretability research including sparse autoencoders and geometric memory structures.

Keywords

AGI alignment; constitutive grammar; human-AI covenant; wisdom traditions; māyā; bilateral corruption

Statements and Declarations

Funding

This research received no external funding. The work was conducted independently by the author.

Competing Interests

The author has no competing interests to declare that are relevant to the content of this article. The author has no financial relationship with any AI company. The use of AI assistance for collaborative development was based solely on assessment of capability for the task.

Ethics Approval

This theoretical and philosophical research did not involve human subjects and did not require ethics approval. The research involves no empirical studies, clinical trials, or collection of personal data.

Consent to Participate

Not applicable. This theoretical research did not involve human participants.

Consent to Publish

Not applicable.

Author Contributions

The sole author is responsible for all aspects of this work, including conception, design, literature review, analysis, framework development, Charter drafting, cross-traditional validation, empirical stress-testing, and manuscript preparation.

Data Availability

All primary sources are cited and publicly available. The Charter document, supplementary materials, and cross-reference tables are available as Online Resources 1–63 accompanying this manuscript. The industry documents reviewed in the manuscript are publicly available from the respective organizations.

AI Assistance Disclosure

Per Springer Nature policy: this manuscript was developed by the human author. An AI language model was used as a tool under direct human oversight to assist with literature synthesis, drafting, structural refinement, and the generation of preliminary coding functions, which were subsequently reviewed, tested, and finalized by the human author. All intellectual contributions, interpretive judgments, and responsibility for the manuscript and its code rest solely with the human author. The AI system did not function as an author.

Acknowledgments

The author acknowledges the use of publicly available AI language models as collaborative tools during the development of this work. The author thanks the open-source communities maintaining the interpretability research tools referenced in this manuscript.